

ECE 463  
Introduction to Computer Networks




Lecture: LAN Interconnects  
Sanjay Rao

# Ethernet

- Chief Advantages:
  - Easy to manage.
  - Cheap
- Concerns:
  - Does not work well under high load
  - Cannot work with too many nodes in system
  - Limitation on the physical distance between the nodes
- How do we get entire Purdue campus connected?


# Ethernet Interconnects

- Repeaters/Hubs
  - Dumb physical layer device that forwards digital signals
  - Devices on either side in the same collision domain
- Bridges/Switches:
  - More intelligent devices that forward data only to hosts needing them.



# Building Larger LANs: Bridges

- Bridges connect multiple IEEE 802 LANs at layer 2.
  - Only forward packets to the right port
  - Reduce collision domain compared with single LAN



# Transparent Bridges

- Overall design goal: **Complete transparency**
  - “Plug-and-play”
  - Self-configuring without hardware or software changes
  - Bridges should not impact operation of existing LANs
- Three parts to transparent bridges:
  - (1) Forwarding of Frames
  - (2) Learning of Addresses
  - (3) Spanning Tree Algorithm

# Frame Forwarding

- Each bridge maintains a **forwarding database** with entries  
< **MAC address**, **port**, **age**>

<b>MAC address:</b>	host name or group address
<b>port:</b>	port number of bridge
<b>age:</b>	aging time of entry


with interpretation:

- a machine with **MAC address** lies in direction of the **port** number from the bridge. The entry is **age** time units old.

# Frame Forwarding 2

- Assume a MAC frame arrives on port x.

**Search if MAC address of destination is listed for ports A, B, or C.**



**Found?**

**Not found ?**

**Forward the frame on the appropriate port**


**Flood the frame, i.e., send the frame on all ports except port x.**



# Address Learning

- In principle, the forwarding database could be set statically (=static routing)
- In the 802.1 bridge, the process is made automatic with a simple heuristic:


The source field of a frame that arrives on a port tells which hosts are reachable from this port.



# Address Learning 2

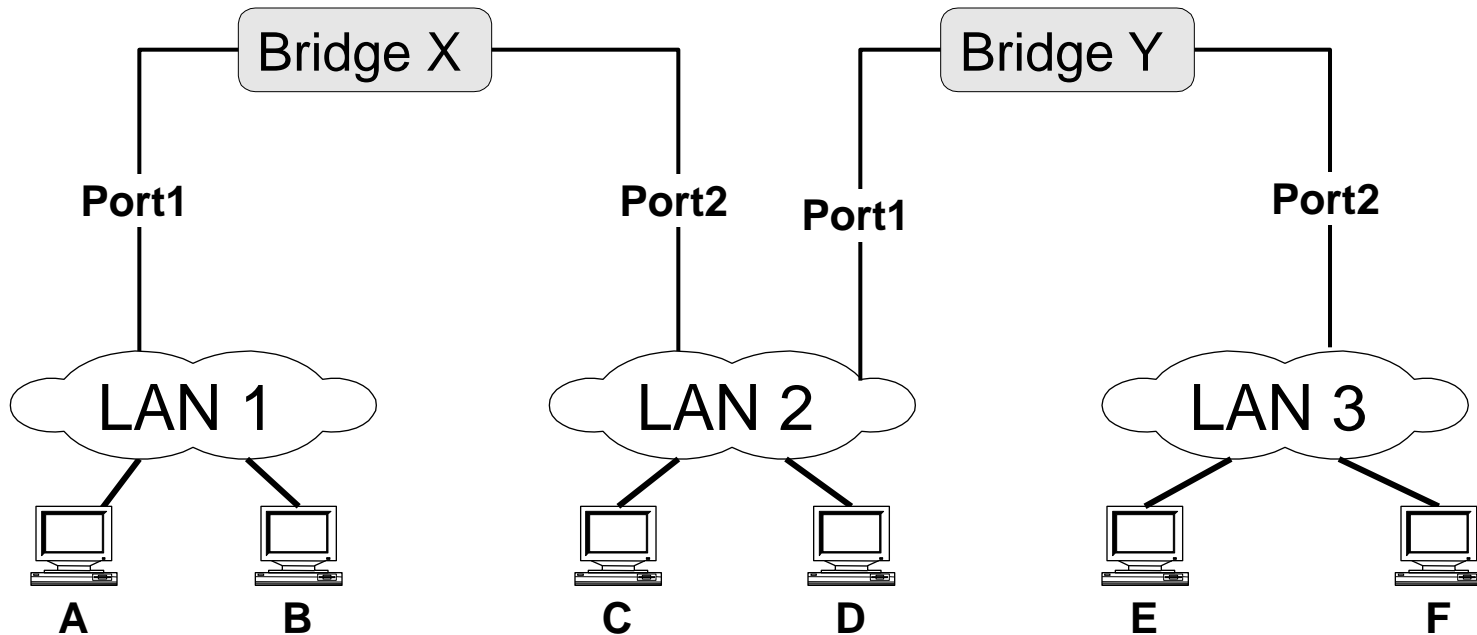
## Algorithm:

- For each frame received, the source stores the source field in the forwarding database together with the port where the frame was received.
- All entries are deleted after some time



# Example

- Consider the following packets:  
<Src=A, Dest=F>, <Src=C, Dest=A>, <Src=E, Dest=C>
- What have the bridges learned?




# Issue

- Consider the two LANs that are connected by two bridges.
- Assume *host n* is transmitting a frame *F* with unknown destination.

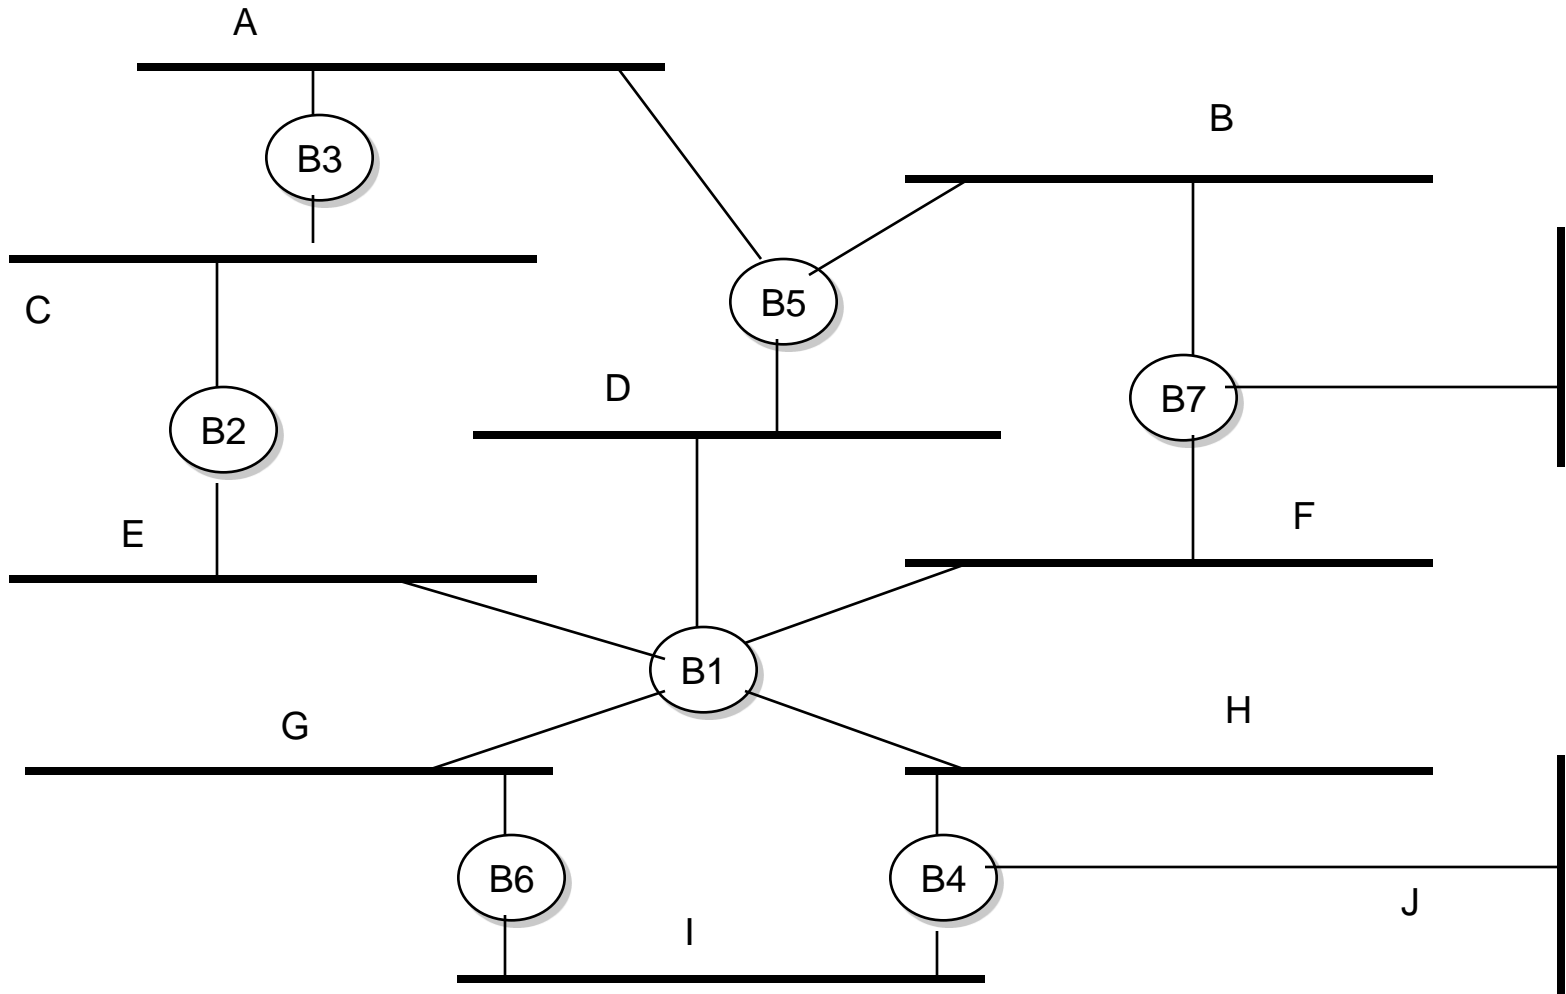
## What is happening?

- Bridges A and B flood the frame to LAN 2.
- Bridge B sees *F* on LAN 2 (with unknown destination), and copies the frame back to LAN 1
- Bridge A does the same.
- The copying continues

Where's the problem? What's the solution ?



# A Bridged Network



# Spanning Trees

- The solution to the loop problem is to not have loops in the topology
- IEEE 802.1 has an algorithm that builds and maintains a **spanning tree** in a dynamic environment.
- Bridges exchange messages to configure the bridge (**Configuration Bridge Protocol Data Unit**, Configuration BPDUs) to build the tree.

# What do the BPDUs do?

With the help of the BPDUs, bridges can:

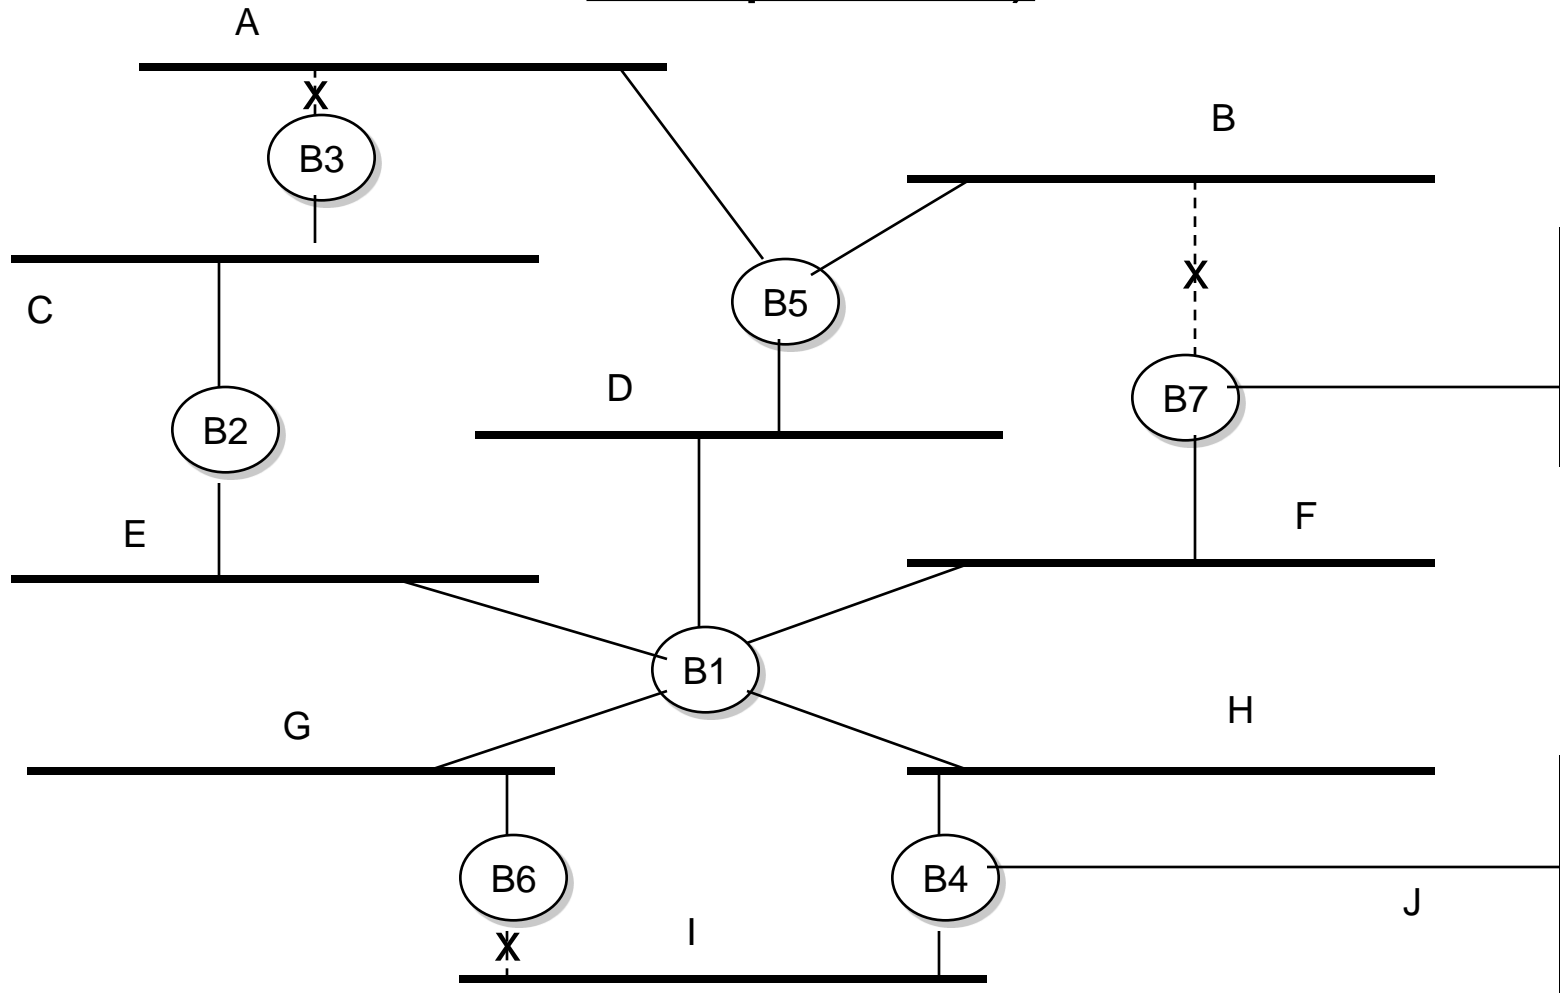
- Elect a single bridge as the **root bridge**.
- Calculate the distance of the shortest path to the root bridge
- Each LAN can determine a **designated bridge**, which is the bridge closest to the root. The designated bridge will forward packets towards the root bridge.
- Each bridge can determine a **root port**, the port that gives the best path to the root.
- Select ports to be included in the spanning tree.

# Concepts

- Each bridge has a unique identifier: **Bridge ID**
- Each port within a bridge has a unique identifier (**port ID**).
- **Root Bridge:** The bridge with the lowest identifier is the root of the spanning tree.
- **Root Port:** Each bridge has a root port which identifies the next hop from a bridge to the root.
- **Root Path Cost:** For each bridge, the cost of the min-cost path to the root, measured in #Hops to the root.
- **Designated Bridge**
  - Bridge on LAN that provides minimal cost path to root for this LAN
  - if two bridges have the same cost, select one with lower identifier
- **Designated Port:**
  - If a bridge is the designated bridge for a LAN, the appropriate port is the designated port.
- Bridge disables all ports which are neither root ports, nor designated ports.



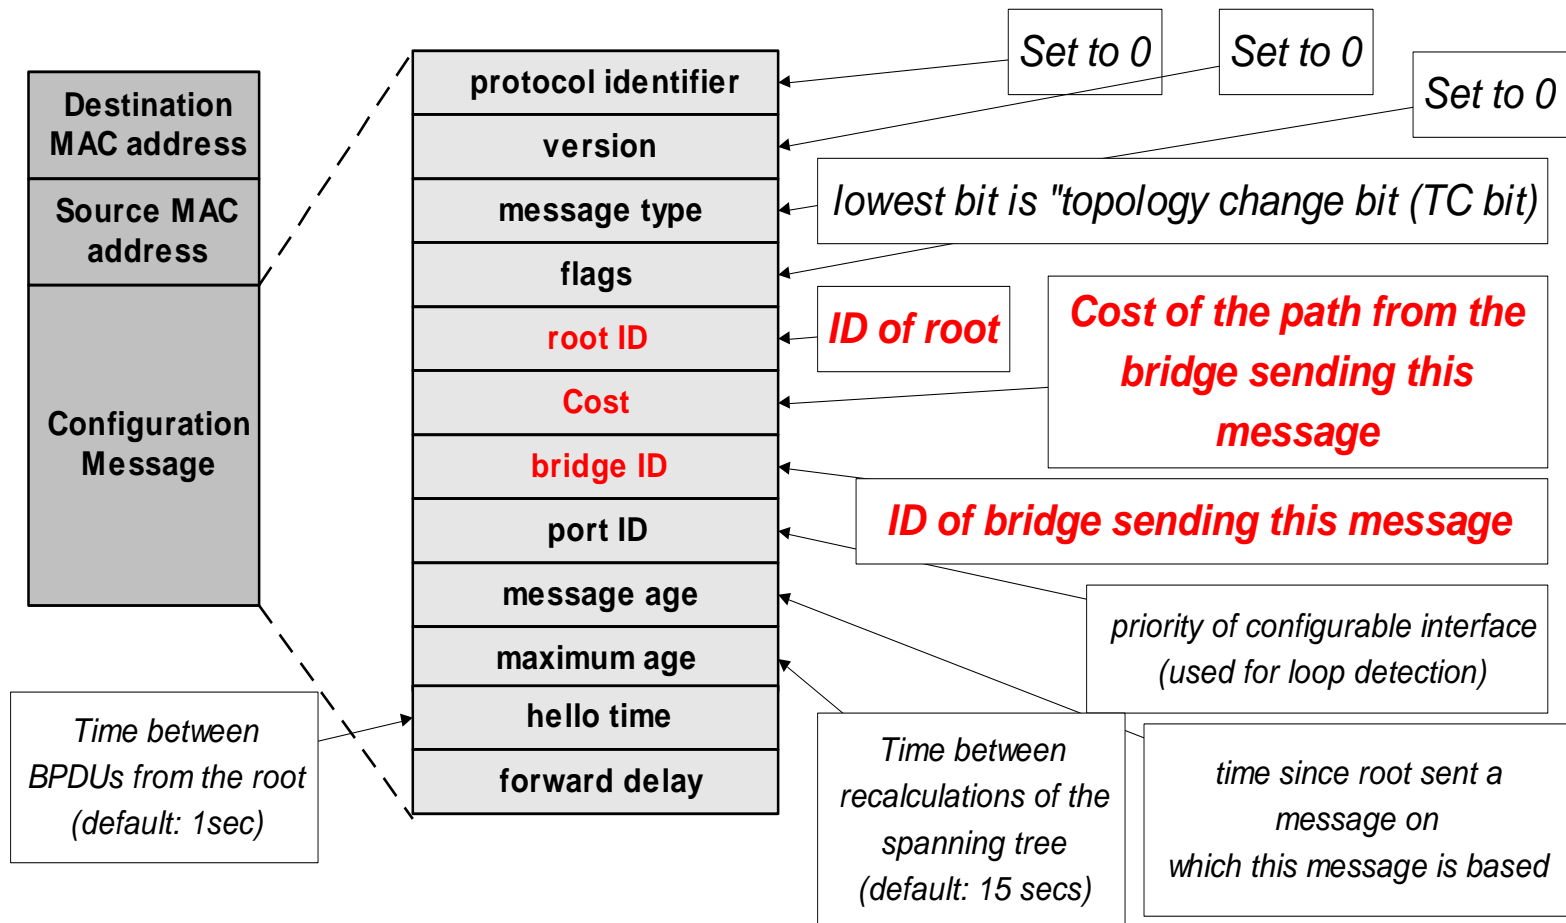
# A Bridged Network (End of Spanning Tree Computation)



# Centralized Vs. Distributed

- Easy to see how this algorithm works if we have the full picture.
- But in reality, each bridge has very limited information
- Need for Distributed algorithms
  - Bridges exchange messages with each other
  - Messages enable them to learn information needed to make the right decisions.
- Topology may not be correct initially, eventually gets to the right one.


# Configuration BPDUs



# Steps of Spanning Tree Algorithm

1. Determine the root bridge
2. Determine the root port on all other bridges
3. Determine the designated port on each LAN


- Each bridge is sending out BPDUs that contain the following information:



root bridge (what the sender thinks it is)  
root path cost for sending bridge  
Identifies sending bridge

# Initialization and Operation


- Initially, all bridges assume they are the root bridge.
- Each bridge B sends BPDUs of this form on all its ports:



- Upon receiving BPDU on a port
  - Bridge checks if received BPDU is “better” than best recorded for that port (initially its own BPDU is best).
  - If “better”, it discards old information for that port, and saves new information, adding 1 to the distance field.

# Ordering of Messages

- We can order BPDUs with the following ordering relation “ $\nwarrow$ ” (let’s call it “lower cost” or “better”):



If ( $R1 < R2$ )

**$M1 \nwarrow M2$**

elseif ( $(R1 == R2)$  and  $(C1 < C2)$ )

**$M1 \nwarrow M2$**

elseif ( $(R1 == R2)$  and  $(C1 == C2)$  and  $(B1 < B2)$ )


**$M1 \nwarrow M2$**

else

**$M2 \nwarrow M1$**

# Determining Root Bridge, Root Port

- If bridge receives message from lower bridge id:
  - It knows it is not the root.
  - It updates its belief of who the root is, say R.
- Bridge B determines the Root Path Cost (Cost) as follows:
  - *If  $B = R$ :* Cost = 0.
  - *If  $B \neq R$ :* Cost = {Smallest Cost in any of BPDUs that were received from R} + 1
- **B's root port** is the port from which B received the lowest cost path to R (in terms of relation " $\nwarrow$ ").
- Knowing R and Cost, B can generate its BPDU (but will not necessarily send it out):





# Calculate the Root Path Cost

## Determine the Root Port

- At this time: B has generated its BPDUs

R	Cost	B
---	------	---

- B will send this BPDUs on one of its ports, say **port x**, only if its BPDUs is lower (via relation “”) than any BPDUs that B received from port x.
- In this case, B also assumes that it is the **designated bridge** for the LAN to which the port connects.





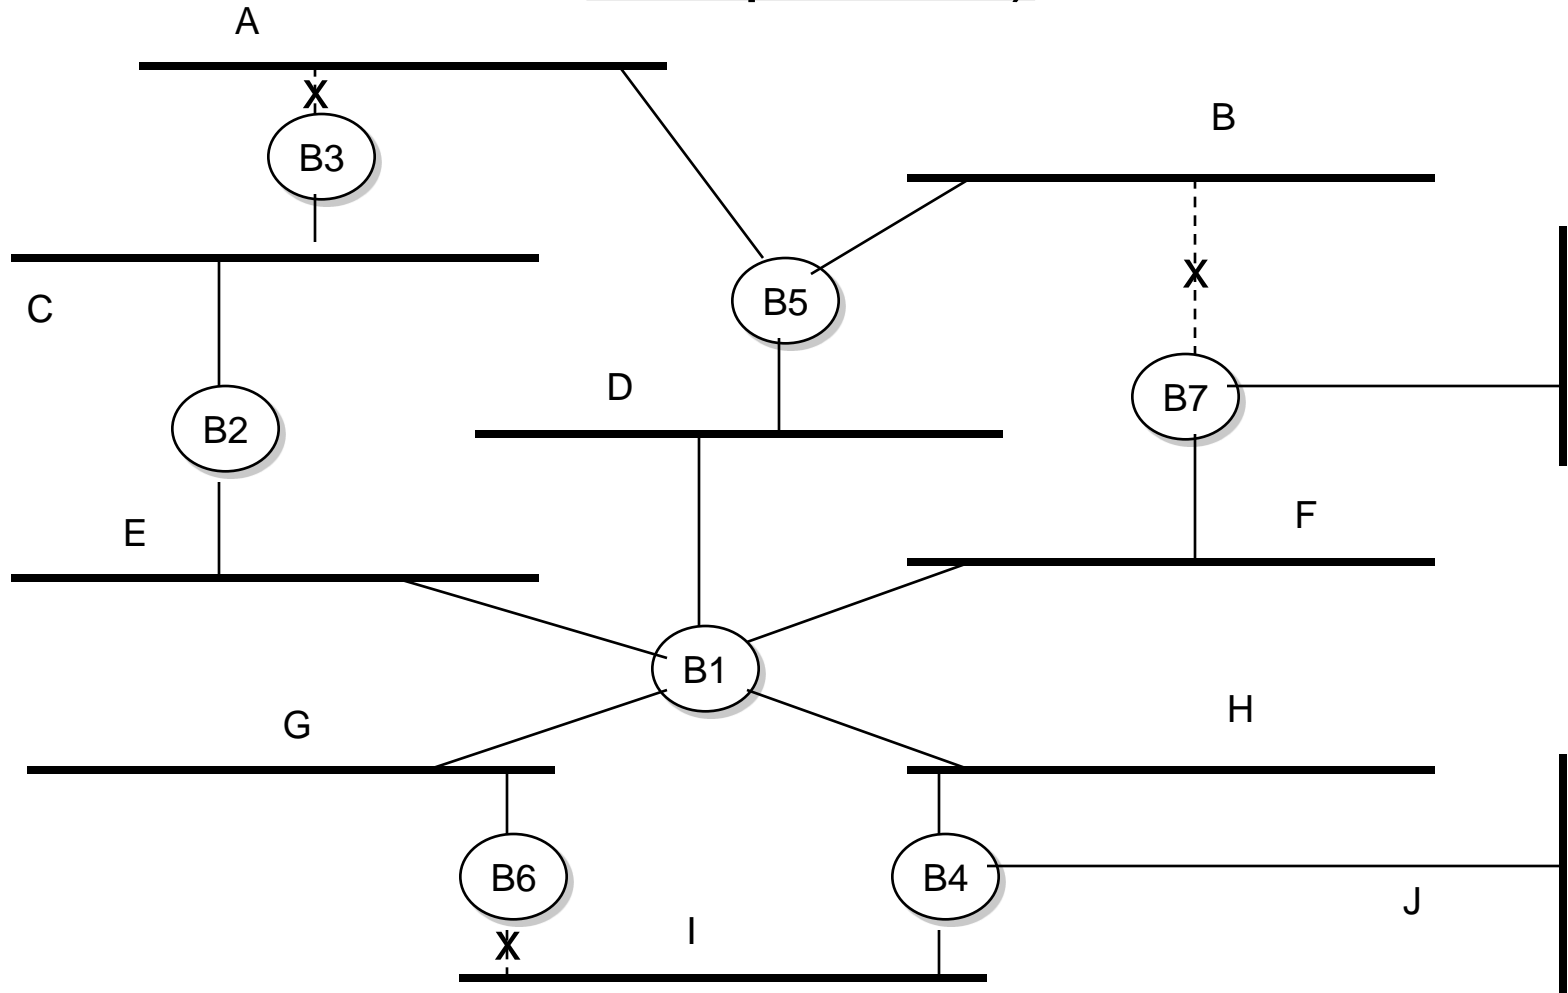
# Selecting the Ports for the Spanning Tree

- At this time: Bridge B has calculated the root, the root path cost, and the designated bridge for each LAN.
- Now **B can decide which ports are in the spanning tree**:
  - B's root port is part of the spanning tree
  - All ports for which B is the designated bridge are part of the spanning tree.
- B's ports that are in the spanning tree will forward packets **(=forwarding state)**
- B's ports that are not in the spanning tree will not forward packets **(=blocking state)**

# Example

- B3 rcvs (B2,0,B2). What happens?
- Makes B2 root. Send what to B5?
  - Sends (B2,1,B3) to B5
- Meanwhile B2 accepts B1 as root
  - Sends (B1,1,B2) to B3
- B3 updates its root to: B1
- Meanwhile B5 accepts B1 as root
  - Sends (B1,1,B5) to B3
- B3: no change to value of root

# A Bridged Network (End of Spanning Tree Computation)



# Ethernet Switches

- Bridges make it possible to increase LAN capacity.
  - Packets are no longer broadcasted - they are only forwarded on selected links
  - Adds a switching flavor to the broadcast LAN
- Ethernet switch is a special case of a bridge: each bridge port is connected to a single host.
  - Simplifies the protocol and hardware used (only two stations on the link) – no longer full CSMA/CD

# Can the Internet be One Big Switched Ethernet?

- Inefficient
  - Too much flooding
- Explosion of forwarding table
  - Need to have one entry for every Ethernet address in the world!
- Poor performance
  - Tree topology does not have good load balancing properties
  - Hot spots
- Etc...

