

# PThreads

Thanks to Professor Sam Midkiff for providing his ECE 563 slides from which these slides are derived.

<https://engineering.purdue.edu/~smidkiff/ece563/slides/PThreads.pdf>

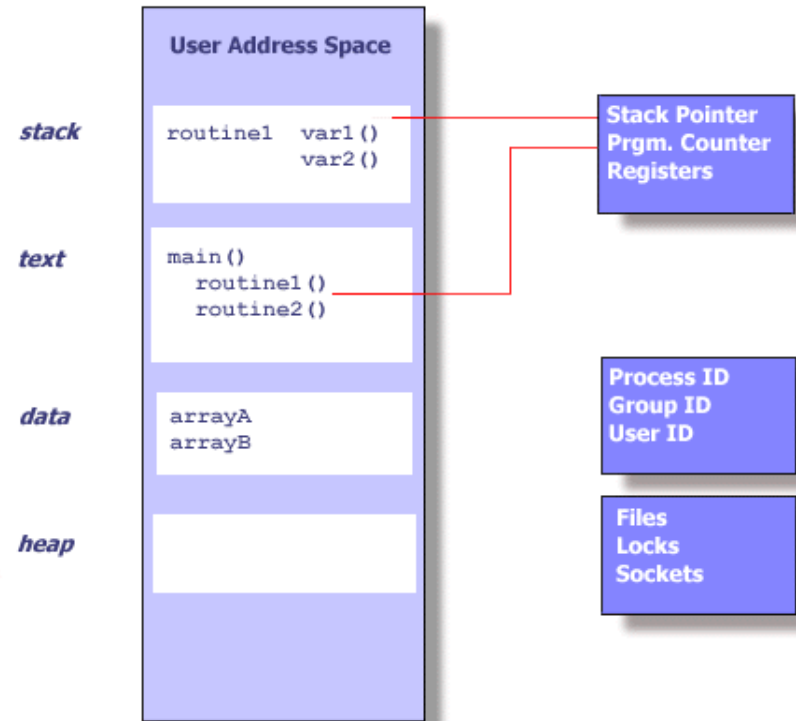
Professor Midkiff's slides were derived from the LLNL tutorial. For a great resource on pthreads see: <https://computing.llnl.gov/tutorials/pthreads/>

# Processes and threads

- Understanding what a thread means knowing the relationship between a process and a thread. A process is created by the operating system.


– Processes contain information about program resources and program execution state, including:

- Process ID, process group ID, user ID, and group ID, address space
- Environment, working directory
- Program instructions, registers, stack, heap
- File descriptors, inter-process communication tools (such as message queues, pipes, semaphores, or shared memory), signal actions
- Shared libraries



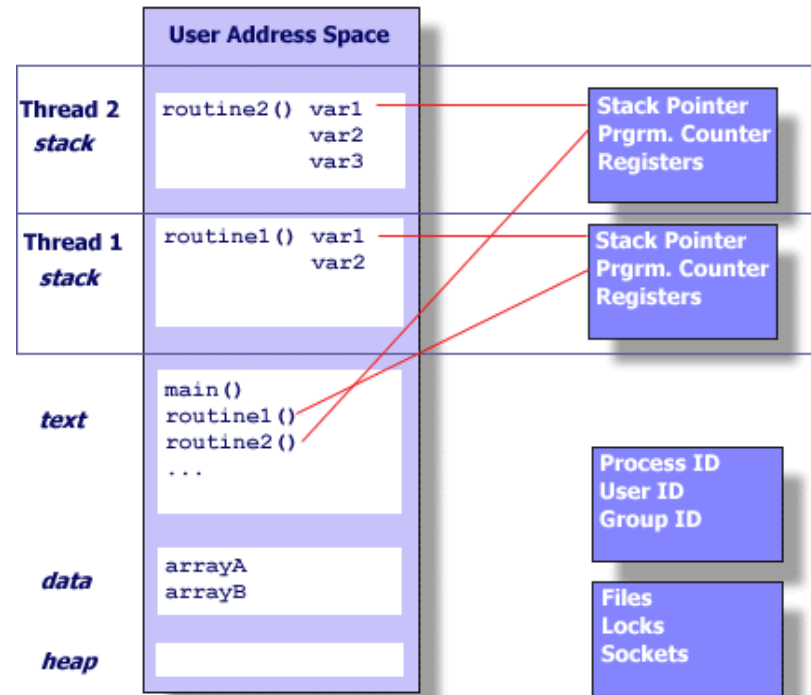
# Processes and threads, cont.

- Threads use and exist within these process resources, yet are able to be scheduled by the operating system and run as independent entities within a process



# Processes and threads, cont.

- A thread can possess an independent flow of control and be schedulable because it maintains its own:
  - Stack pointer
  - Registers
  - Scheduling properties (such as policy or priority)
  - Set of pending and blocked signals
  - Thread specific data.



# Processes and threads, cont.

- A process can have multiple threads, all of which share the resources within a process and all of which execute within the same address space
- Within a multi-threaded program, there are at any time multiple points of execution

# Processes and threads, cont.

- Because threads within the same process share resources:
  - Changes made by one thread to shared system resources (such as closing a file) will be seen by all other threads
  - Two pointers having the same value point to the same data
  - Reading and writing to the same memory locations is possible, and therefore requires explicit synchronization by the programmer

# What are Pthreads?

- Historically, hardware vendors implemented their own proprietary versions of threads.
  - Standardization required for portable multi-threaded programming
  - For Unix, this interface specified by the IEEE POSIX 1003.1c standard (1995).
    - Implementations of this standard are called POSIX threads, or Pthreads.
    - Most hardware vendors now offer Pthreads in addition to their proprietary API's
    - Pthreads are defined as a set of C language programming types and procedure calls, implemented with a pthread.h header/include file and a thread library

# Creating threads

- [pthread\\_create](#)(thread, attr, start\_routine, arg)
- This routine creates a new thread and makes it executable. Typically, threads are first created from within `main()` inside a single process.
  - *start\_routine* is the C routine that the thread will execute once it is created. A single argument may be passed to *start\_routine* via *arg* as a void pointer.
  - The *attr* parameter is used to set thread attributes. Can be an object, or NULL for the default values. (For Lab2, you only need to set to NULL)
  - The maximum number of threads that may be created by a process is implementation dependent.



# Terminating threads

- Multiple methods in general. For our purposes:
  - **Recommended for Lab 2:** The thread returns from its starting routine (the main routine for the initial thread)
  - Another option: The thread makes a call to the `pthread_exit` subroutine
    - Returning from starting routine is equivalent to calling `pthread_exit` with the value supplied in the return statement

- pthread\_join(threadId, status)
- The pthread\_join() subroutine blocks the calling thread until the specified *threadId* thread terminates
- The programmer is able to obtain the target thread's termination return status if specified in: pthread\_exit(), in the *status* parameter
  - Or value when returning from the thread's starting\_routine

# Protecting access to memory locations shared by threads

Example code:

Global variable:

“balance”: how much money left in a bank account

Two threads:

- Thread 1 is withdrawing money
- Thread 2 is depositing money

Operations occur concurrently.

Both threads access the global variable “balance”

Need to be careful to protect accesses to the variable.

# Protecting access to memory locations shared by threads

Global variable: “balance”: how much money left in a bank account

Two separate threads execute withdrawalFn, and depositFn

```
void * withdrawalFn(void * arg){
    int curr_balance = balance;
    curr_balance = curr_balance - withdrawal_amount;
    sleep(1);
    balance = curr_balance;

    return NULL;
}

void * depositFn(void * arg){
    int curr_balance = balance;
    curr_balance = curr_balance + deposit_amount;
    sleep(1);
    balance = curr_balance;

    return NULL;
}
```

# Example race condition that can occur

Thread 1 (deposit)	Thread 2 (withdraw)	Balance
Read balance: \$1000		\$1000
	Read balance: \$1000	\$1000
	Withdraw \$200	\$1000
Deposit \$200		\$1000
Update balance \$1000+\$200		\$1200
	Update balance \$1000-\$200	\$800

# Using locks

- pthread\_mutex\_lock(mutex)
  - Acquire lock if available
  - Otherwise wait until lock is available
- pthread\_mutex\_unlock(mutex)
  - Release the lock to be acquired by another pthread\_mutex\_lock call
  - Cannot make assumptions about which thread acquire the lock next

# Protecting access to shared variables using locks

```
void * withdrawalFnLock(void * arg){
    pthread_mutex_lock(&lock);
    int curr_balance = balance;
    curr_balance = curr_balance - withdrawal_amount;
    sleep(1);
    balance = curr_balance;
    pthread_mutex_unlock(&lock);

    return NULL;
}

void * depositFnLock(void * arg){
    pthread_mutex_lock(&lock);
    int curr_balance = balance;
    curr_balance = curr_balance + deposit_amount;
    sleep(1);
    balance = curr_balance;
    pthread_mutex_unlock(&lock);

    return NULL;
}
```

# Example execution with locks

Thread 1 (deposit)	Thread 2 (withdraw)	Balance
	GetLock()	\$1000
	<b>Read balance: \$1000</b>	\$1000
	<b>Withdraw \$200</b>	\$1000
	<b>Update balance \$1000-\$200</b>	\$800
	ReleaseLock()	\$800
GetLock()		\$800
<b>Read balance: \$800</b>		\$800
Deposit \$200		\$800
<b>Update balance \$800+\$200</b>		\$1000
ReleaseLock()		\$1000



# Creating threads

- [pthread\\_create](#)(thread, attr, start\_routine, arg)
- This routine creates a new thread and makes it executable. Typically, threads are first created from within `main()` inside a single process.
  - *start\_routine* is the C routine that the thread will execute once it is created. A single argument may be passed to *start\_routine* via *arg* as a void pointer.
  - The *attr* parameter is used to set thread attributes. Can be an object, or NULL for the default values. (For Lab2, you only need to set to NULL)
  - The maximum number of threads that may be created by a process is implementation dependent.

time

Thread 0

Thread  $k$

$t = 0;$

`pthread_create(..., f, t);`

$t = 1$

`pthread_create(..., f, t);`

$t = 2$

*thread spawn*

`f(t);`

`x = t;`

What is the value of  $t$  that is used in this call to `f`?  
The value is indeterminate.

# Passing arguments to a thread

- Thread startup is non-deterministic
- It is implementation dependent
- If we do not know when a thread will start, how do we pass data to the thread knowing it will have the right value at startup time?
  - Don't pass data as arguments that can be changed by another thread
  - In general, use a separate instance of a data structure for each thread.

# Passing data to a thread (a simple integer)

```
int *taskids[NUM_THREADS];
for(t=0;t < NUM_THREADS;t++) {
    taskids[t] = (int *)
                    malloc(sizeof(int));
    *taskids[t] = t;
    printf("Creating thread %d\n", t);
    rc = pthread_create(&threads[t], NULL,
                        PrintHello,
                        (void *) &t);
    ...
}
```