

## Process Synchronization

ECE595

Aug 30

Y. Charlie Hu



1

## Review of Lec2



- What is a process?
- Benefits of process isolation?
- Process state transition diagram
- What is PCB? What are stored in it?
  
- Why concurrent processes?
- Process creation and termination
  - `fork()`
  - `exec()`

2

## Lec1: Polycs vs. Mechanisms

When `fork()` creates a child process



- Execution possibilities?
  - Parent and child execute concurrently
  - Parent waits till child finishes
    - Concurrency?
- Address space possibilities?
  - Child duplicates that of parent
  - Child has a program loaded into it

3


## [lec3] Process Creation -- UNIX examples



- `fork()` system call creates a duplicate of the original process
  - Should have been called “clone()”
  - Allows the two to communicate (one time)
  - How to disambiguate who is who?
- `exec()` system call used after a `fork` to replace the process' code/address space with a new program


4

## fork()



5

## exec("b.out")



Afterwards, only one thing about the process was kept, which is?

6

## Roadmap

- What is a process?
- Concurrent processes
- Process creation
- Non-preemptive CPU scheduling
- Process synchronization

7

## Multiprogramming needs CPU scheduling

- Without any hardware support, what can the OS do to a running process?

8

## System calls may trigger Scheduler

- Block – wait on some event/resource
  - Network packet arrival (e.g., `recv()`)
  - Keyboard, mouse input (e.g., `getchar()`)
  - Disk activity completion (e.g., `read()`)
- Yield – give up running for now

9

## Non-Preemptive Scheduler


- A non-preemptive scheduler: a scheduler that is only invoked by explicit block/yield calls, or terminations
  - Only method when there is no timer!
- The simplest form
 

**Scheduler:**

  - save current process state (into PCB)
  - choose next process to run
  - dispatch (load PCB and run)
- Used in Windows 3.1, Mac OS

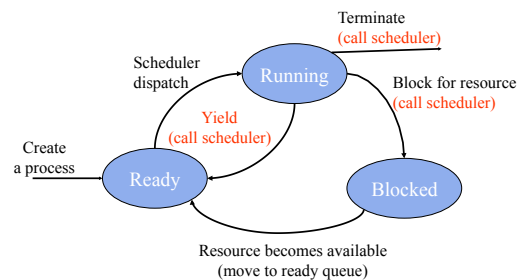
10

## Our Friend -- the Transition Diagram



11

## Process State Transition of Non-Preemptive Scheduling



How does a process terminate itself?

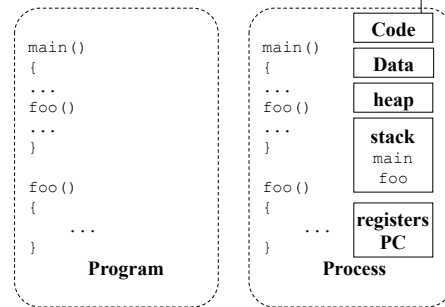
12

## Context Switch

- Definition:  
switching the CPU to run another process, which involves (1) saving the state of the old process and (2) loading the state of the new process
- What state?

13

## [lec3] Program vs. Process



14

## Context Switch

- Definition:  
switching the CPU to run another process, which involves (1) saving the state of the old process and (2) loading the state of the new process
- What state?
  - What about L1/L2 cache content?

15

## Context Switch

- Definition:  
switching the CPU to another process, which involves saving the state of the old process and load the state of the new process
- What state?
- Where to store them?

16

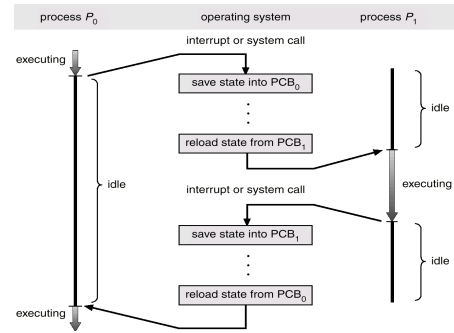
## [lec3] Process Control Block (Process Table)

- Process management info
  - State (ready, running, blocked)
  - PC & Registers, parents, etc
  - CPU scheduling info (priorities, etc.)
- Memory management info
  - Segments, page table, stats, etc
- I/O and file management
  - Communication ports, directories, file descriptors, etc.



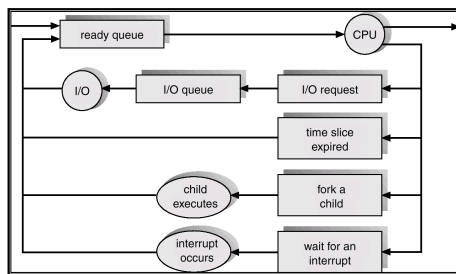
17

## Context Switch



18

## Process Scheduling Diagram



19

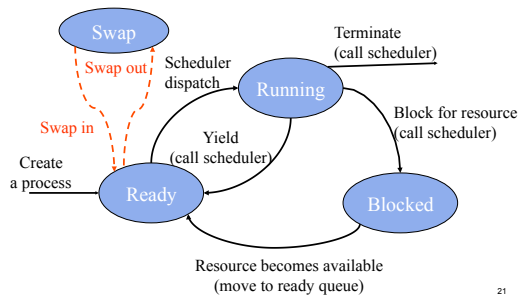
## Physical Memory & Multiprogramming

- Want to run many programs
- Programs need memory to run
- Memory is a scarce resource
- What happens when  $M(a) + M(b) + M(c) > \text{physical mem?}$



20

## Add Job Swapping to State Transition Diagram



## Quiz

Can you swap out from Blocked?

22

## Process synchronization

- Cooperating processes may share data via
  - shared address space (code, data, heap) by using threads
  - shared memory objects (used in lab2)
  - Files
  - (Sending messages)
- What can happen if processes try to access shared data (address) concurrently?
  - Sharing bank account with sibling:
    - At 3pm: If (balance > \$10) withdraw \$10
- How hard is the solution?

23

## “Got Milk?”

- “Too much milk” problem

24

## Process synchronization needs help from OS!



25

## Mutual exclusion & Critical Section



- **Critical section** – a section of code, or collection of operations, in which only one process shall be executing at a given time
- **Mutual exclusion** - mechanisms that ensure that only one person or process is doing certain things at one time (others are excluded)

26

## Example of Critical Section



- Concurrent accesses to shared variables, at least one of which is write

P0:	P1:	P2:
Read note;	Read note;	write note;
...	...	...

27

## Desirable properties of MuEx



- **Fair**: if several processes are waiting, let each in eventually
- **Efficient**: don't use up substantial amounts of resources when waiting (e.g. no busy waiting)
- **Simple**: should be easy to use (e.g. just bracket the critical sections)

28

## Desirable properties of processes using MuEx

- Always lock before manipulating shared data
- Always unlock after manipulating shared data
- Do not lock again if already locked
- Do not unlock if not locked by you
- Do not spend large amounts of time in critical section

29

## Mutual Exclusion provided by OS (and language/compiler)

- Locks
  - Alone can solve simple problems
- More advanced
  - Semaphore
  - Lock and condition variable
    - Lock alone is not flexible enough
  - Monitor
- Each primitive itself is atomic!

30

## Lock (aka mutex)

Init: lock = 1; // 0 means held; 1 means free

```
lock_acquire(lock)      lock_release(lock)

{
    while (lock==0);
    lock--;
}

{
    if (lock == 0)
        lock++;
}
```

- Each primitive is atomic
- In reality, lock is not implemented as above!
  - The waiting process is put to sleep

31

## “Too much milk” problem with locks

```
Acquire(lock);
if (noMilk)
    buy milk;
Release(lock); } Critical Section
```

- What is the problem with this solution?

32



## Deep thinking

- How can we solve the problem?

```
if ( noMilk ) {
    if (noNote) {
        leave note;
        buy milk;
        remove note;
    }
}

if ( noMilk ) {
    if (noNote) {
        leave note;
        buy milk;
        remove note;
    }
}
```

33

## Often times, we have to wait for shared resources

- In “too much milk”, we just needed to check
- Often times, before accessing shared resources, we have to wait ...

35

## Producer & Consumer Problem (1-pool version)

- **Producer**: creates copies of a resource
- **Consumer**: uses up (destroys) copies of a resource. (may produce something else)
- **Buffer**: used to hold resource produced by producer before consumed by consumer
- **Synchronization**: keeping producer & consumer in sync
- Happens inside OS all the time (e.g. I/Os)
  - How about in real life?

36

## Producer & Consumer – solution using locks?

Producer	Consumer
<pre>while (1) {</pre>	<pre>While (1) {</pre>
<pre>    produce an item;</pre>	<pre>    while (buffer is empty);</pre>
<pre>    while (buffer is full);</pre>	<pre>    remove an item;</pre>
<pre>    insert item into buffer</pre>	<pre>    consume the item</pre>
<pre>}</pre>	<pre>}</pre>

37

## Often times, we have to wait for shared resources



- Busy waiting is a bad idea
- Checking resources itself needs to be in critical section
- Busying waiting inside CS even worse!
  - No one else can check!

→ Need a more powerful sync. primitive!

42

## Reading assignment



- Chapter 6 (process synchronization)
  - Read up materials covered

43