

Communication with Messages, Wait-Free Synchronization

ECE595
Sep 11
Y. Charlie Hu



1

Classic Synchronization Problems



1. Producer-consumer problem (bounded buffer problem)
2. Readers-writers problem
3. Dining philosophers problem

3

Readers-Writers problem



Abstraction of concurrent access to shared data problem

- A data object is shared among multiple processes

Reader:

```
While (1) {
```

```
  acq(mutex)
```

```
  read();
```

```
  rel(mutex)
```

```
}
```

Writer:

```
While (1){
```

```
  acq(mutex)
```

```
  write();
```

```
  rel(mutex) /* abstraction */
```

```
}
```

4

Readers-Writers problem



Abstraction of concurrent access problem

- A data object is shared among multiple processes
- **Allow concurrent reads, but exclusive writes**
 - Implication: need to move read() and write() outside Critical Sec
 - Can we use semaphore to count readers/writers?

Reader:

```
acq(mutex)
```

```
????
```

```
rel(mutex)
```

```
read();
```

```
acq(mutex)
```

```
???
```

```
rel(mutex)
```

Writer:

```
acq(mutex)
```

```
????
```

```
rel(mutex)
```

```
write();
```

```
acq(mutex)
```

```
???
```

```
rel(mutex)
```

5

Readers-Writers problem

Abstraction of concurrent access problem

- A data object is shared among multiple processes
- Allow concurrent reads, but exclusive writes
- Solution needs lock, counting, and semaphores!
- Constraints:
 - Writers can only proceed if there are no active readers/writers
→ use semaphore `OKtoWrite`
 - Readers can proceed only if there are no active/waiting writers
→ use semaphore `OKtoRead`
 - To keep track of how many are reading / writing / waiting
→ use some shared variables, called *state variables*
 - Only one process manipulates state variable at once
→ use a lock `Mutex`

6

Readers-Writers problem (cont)

- State variables:
 - `AR` = number of active readers
 - `WR` = number of waiting readers
 - `AW` = number of active writers
 - `WW` = number of waiting writers

`AW` is always 0 or 1
`AR` and `AW` can not both be non-zero
- Initialization:
 - `OKtoRead` = 0;
 - `OKtoWrite` = 0;
 - `Mutex` = 1;
 - `AR` = `WR` = `AW` = `WW` = 0;
- Scheduling: writers get preference

7

Readers-Writers problem (cont)

• Reader

```
wait(mutex)
if ((AW + WW) == 0) {
    signal(OKtoRead);
    AR++;
} else {
    WR++;
}
signal(mutex)
```

```
wait(OKtoRead);
read necessary data;
```

```
wait(mutex);
AR--;
if (AR == 0 && WW > 0) {
    /* no need to check AW == 0 */
    signal(OKtoWrite);
    AW++;
    WW--;
}
signal(mutex);
```

8

Readers-Writers problem (cont)

• Writer

```
wait(mutex)
if ((AW + AR + WW) == 0) {
    signal(OKtoWrite);
    AW++;
} else {
    WW++;
}
signal(mutex);
```

```
wait(OKtoWrite);
write necessary data;
```

```
wait(mutex);
AW--;
if (WW > 0) {
    signal(OKtoWrite);
    AW++;
    WW--;
} else while (WR > 0) {
    signal(OKtoRead);
    AR++;
    WR--;
}
signal(mutex);
```

9

Dining philosophers problem

Abstraction of concurrency-control problems

The need to allocate several resources among several processes while being deadlock-free and starvation-free



Roadmap

- Interprocess communication *with shared data*
 - Synchronization with locks, semaphores, condition var
 - Classic sync. problem 1: producer-consumer
 - Semaphore implementations (uniprocessor, multiprocessor)

Today:

- Classic sync. problems 2 & 3
- Interprocess communication *with messages*
- Wait-free synchronization


11

Inter-process Communication with Messages

- Messages provide for communication *without shared data*
 - One process or the other owns the data, (guaranteed) never two at the same time
 - Think about usmail

12

Big Picture



13

Why use messages?

- Many types of applications fit into the model of processing a sequential flow of information
- Communication across address spaces – no side effects
 - Less error-prone
 - They might have been written by diff programmers who aren't familiar with code
 - They might not trust each other
 - They may not be running on different machines!
 - Examples?

14

Message Passing API

- Generic API
 - `send(mailbox, msg)`
 - `recv(mailbox, msg)`
- What is a mailbox?
 - A *buffer* where messages are stored between the time they are sent and the time when they are received
- What should “msg” be?
 - Fixed size msgs
 - Variable sized msgs: need to specify sizes

15

Buffering leads to design options

- When should `send()` return?
- When should `recv()` return?

16

Send

- *Fully Synchronous*
 - Will not return until data is received by the receiving process
- *Synchronous*
 - Will not return until data is received by the mailbox
 - Block on full buffer
- *Asynchronous*
 - Return immediately
 - Completion
 - Require the application to check status (appl polls)
 - Notify the application (OS sends interrupt)
 - Block on full buffer

17

Receive



- *Synchronous*
 - Return data if there is a message
 - Block on empty buffer
- *Asynchronous*
 - Return data if there is a message
 - Return status if there is no message (probe)

18

OS implementation



- What is the conceptual problem for OS implementation here?
 - Assume sender and receiver are on the same machine

19

Buffering



- No buffering
 - Sender must wait until the receiver receives the message
 - Rendezvous on each message
- Bounded buffer
 - Finite size
 - Sender blocks when the buffer is full
 - Receiver blocks when the buffer is empty
 - Using lock/condition variable (or semaphore)

20

Direct Communication



- Each process must name the sending or receiving process
- A communication link
 - is set up between the pair of processes
 - is associated with exactly two processes
 - exactly one link between each pair of processes

```
P: send( process Q, msg )
Q: rcv( process P, msg )
```

21

Producer-Consumer Problem with Message Passing

```
Producer(){
  while (1) {
    ...
    produce item
    ...
    send( consumer, item);
  }
}
```

```
Consumer(){
  while (1) {
    recv( producer, item );
    ...
    consume item
    ...
  }
}
```

22

Indirect Communication

- Use a “mailbox” or “ports” to allow **many-to-many communication**
 - Mailbox typically owned by the OS
 - Requires open/close a mailbox before allowed to use it
- A “link”
 - is set up among processes only if they have a shared mailbox
 - Can be associated with more than two processes

```
P: open (mailbox); send( mailbox, msg);
   close(mailbox)
Q: open (mailbox); recv( mailbox, msg );
   close(mailbox)
```

23

The big debate in parallel computing: Messaging vs. Sharing Data

- Two programming models are equally powerful
- But result in very different-looking programming styles
- Most people find shared-data programming easier to work with
 - Debugging?
- What about machines that do not share memory?
 - Can be simulated in software [SDSM – hot topic in 80-90’ s]
 - But often not as efficient as message passing

24

Roadmap

- Interprocess communication **with shared data**
 - Synchronization with locks, semaphores, condition var
 - Classic sync. problem 1: producer-consumer
 - Semaphore implementations (uniprocessor, multiprocessor)

Today:

- **Classic sync. problems 2 & 3**
- **Interprocess communication with messages**
- **Wait-free synchronization**

25

[lec6] Uniprocessor solution: disable interrupts!

```

void wait(semaphore s)
{
    disable interrupts;
    if (s->count > 0) {
        s->count --;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    enable interrupts;
    sleep(); /* re-dispatch */
}

void signal(semaphore s)
{
    disable interrupts;
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        process = removeFirst(s->q);
        wakeup(process);
        /* put process on Ready Q */
    }
    enable interrupts;
}

```

26

[lec6] Use TAS to implement semaphores on multiprocessor?

```

void wait(semaphore s)
{
    disable interrupts;
    while (tsa(&lock, 1) == 1);
    if (s->count > 0) {
        s->count --;
        lock = 0;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    lock = 0;
    sleep(); /* re-dispatch */
    enable interrupts;
}

void signal(semaphore s)
{
    disable interrupts;
    ???
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(thread);
        /* put process on Ready Q */
    }
    ???
    enable interrupts;
}

```

Do we still need to disable interrupts?

27

Wait-free Synchronization

- Finally we need tsa or ldl&stc anyway to implement sync. primitives (on multiprocessors)
- Can we design data structures in a way that allows safe concurrent accesses?
 - no mutual exclusion necessary
 - no possibility of deadlock
 - only using tsa / ldl&stc
 - no busy waiting


28

Simple example – Queue insertion

```

typedef struct {
    QItem *item;
    QElem *next;
} QElem;


```



29

Simple example – Queue insertion

```
typedef struct {
    QItem *item;
    QElem *next;
} QElem;
```



30

Singly-linked Queue Insertion

```
QElem *queue;

void Insert(item) {
    QElem *new = malloc(sizeof(QElem));
    new->item = item;
    new->next = queue;
    queue = new;
}
```

31

Wait-free Synchronization

- Design data structures in a way that allows safe concurrent accesses
 - no mutual exclusion (lock acquire & release) necessary
 - no possibility of deadlock
- Approach: use a single atomic operation to
 - commit all changes
 - move the shared data structure from one consistent state to another

32

Read-modify-write on CISC

- Most CISC machines provide atomic *read-modify-write* instruction
- Assume a test-and-set instruction

```
X = TAS(addr, old_value, new_value);

read value V at addr;
if (V == old_value) set it to new_value;
return V;
```

33

Singly-linked Queue Insertion

```
QElem *queue;

void Insert(item) {
    QElem *new = malloc(sizeof(QElem));
    new->item = item;
    do {
        new->next = queue;
    } while (tsa(&queue, new->next, new) !=
            new->next);
}
```

Is this
busy
waiting
a problem?

34

Limitation

- Example only works for simple data structures where changes can be committed with *one store instruction*
- What about more complex data structures?

35

More General Approach

- Maintain a pointer to the “master copy” of the data structure
- To modify,
 1. remember current value of the master pointer
 2. copy shared data structure to a scratch location
 3. modify copy
 4. *atomically*:
 - verify that master pointer has not changed
 - write pointer to refer to new master
 5. if verification fails (another process interfered), start over at step 1
- Downside?
 - When does it work reasonably well?

36

Next week

- We will move on to CPU scheduling !

37