

CPU Scheduling

ECE595
Sep 13

Y. Charlie Hu



1

Roadmap So Far



- Processes, (threads), creation
- Inter-process comm by sharing data → process synchronization
 - OS-provided sync. Primitives
 - Mutual exclusion & Critical section
 - Semaphore (binary semaphore)
 - Lock / condition variable
 - Classic sync. Problems
 - Producer-consumer problem
 - Reads-writers problem
 - Dining Philosophers problem (deadlock)
 - Semaphore implementation in OS
 - Wait-free synchronization
- Inter-process comm by messaging (mailboxes)

2

CPU Scheduling



- CPU scheduling is the basis of multiprogrammed operating systems
- By switching the CPU among processes, the OS can make the CPU/computer maximally utilized

3

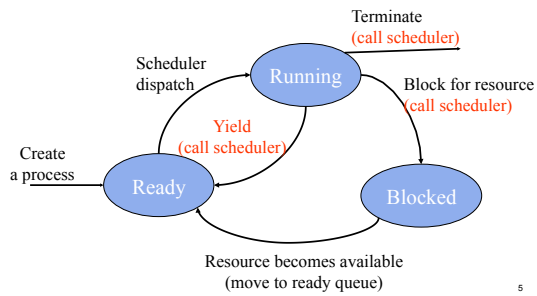
Hardware Support



- Without hardware support, can we do anything other than non-preemptive scheduling?

4

[lec3] Process State Transition of Non-Preemptive Scheduling



5


Timesharing Systems

- **Timesharing** systems support interactive use
 - each user feels he/she has the entire machine
- How?
 - optimize response time
 - based on time-slicing

6

Timer Interrupts

- Using timer interrupt to do CPU management
- Timer interrupt
 - generated by hardware
 - setting requires privilege
 - delivered to the OS



7

Using Interrupts For Scheduling

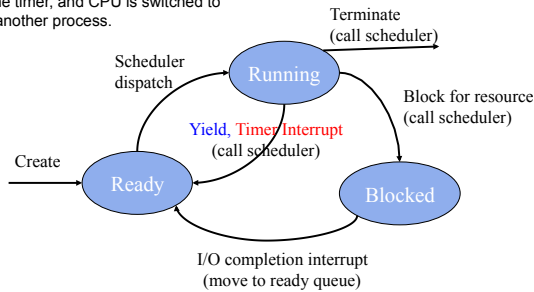
Basic idea

- before moving process to running, OS sets timer
- if process yields/blocks, clear timer, go to scheduler
- If timer expires, go to scheduler

8

Preemptive Scheduling

A running process is interrupted by the timer, and CPU is switched to run another process.



9

[lec4] Context Switch

- **Definition:**
switching the CPU to another process, which involves saving the state of the old process and loading the state of the new process

- What state?
- Where to store them?

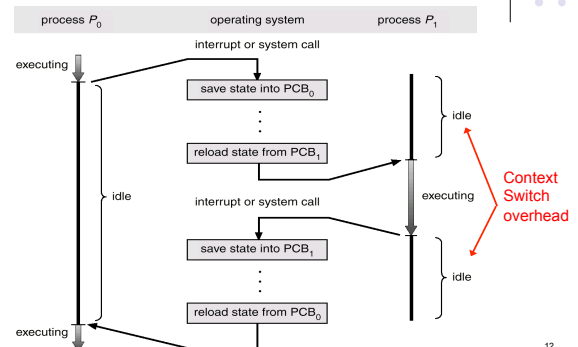
10

Process Control Block (Process Table)

- **Process management info**
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- **Memory management info**
 - Segments, page table, stats, etc
- **I/O and file management**
 - Communication ports, directories, file descriptors, etc.

11

Context Switch



12

Preemptive Scheduling Considerations



- Timer granularity
 - Finer timers = more responsive
 - Coarser timers = more efficient
- CPU Accounting (CPU running stats)
 - Used by the scheduler
 - Useful for the programmer

13

OS as a Resource Manager: Allocation vs. Scheduling



- **Allocation** (spatial)
 - **Who gets what.** Given a set of requests for resources (e.g. memory), which processes should be given which resources (e.g. how much memory & where) for best utilization
- **Scheduling** (temporal)
 - **How long can they keep it.** When more resources (e.g. 10 CPUs) are requested than can be granted (e.g. 1 CPU), in what order can they be serviced?

14

[lec1] Separating Policy from Mechanism



Mechanism – tool to achieve some effect

Policy – decisions on how to use tool
examples:

- All users treated equally
- All program instances treated equally
- Preferred users treated better

Separation leads to flexibility

15

Preemptive CPU Scheduling



- What is in it?
 - **Mechanism + policy**
 - Mechanisms fairly simple
 - Policy choices harder

16

[lec1] Brief History of Computer Systems (1)



- In the beginning, 1 user/program at a time
- Simple **batch** systems were 1st real OS
 - **Spooling and buffering** allowed jobs to be read ahead of time
- **Multiprogramming** systems provided increased utilization (throughput)
 - multiple runnable jobs loaded in memory
 - overlap I/O with computation
 - benefit from asynchronous I/O devices
 - 1st instance where the OS must allocate and schedule resources
 - CPU scheduling
 - Memory management
 - Protection

17

[lec1] Brief History of Computer Systems (2)



- **Timesharing** systems support interactive use
 - Logical extension of multiprogramming
 - optimize response time by frequent time-slicing multiple jobs
 - each user feels he/she has the entire machine
 - permits interactive work
- Most systems today are timesharing (focus of this class)

18

[lec1] Is there a perfect OS? (resource manager, abstract machine)



Fairness
Efficiency

Portability
Interfaces

Security
Robustness

- Conflicting goals
 - Fairness vs efficiency
 - Efficiency vs portability
 - ...
- Furthermore, ...

19

Challenges in Policy



- Flexibility - variability in job types
 - Long vs. short
 - Interactive vs. non-interactive
 - I/O-bound vs. compute-bound
- Issues
 - Short jobs shouldn't suffer
 - (Interactive) Users shouldn't be annoyed

20

Challenges in Policy (cont)

- Fairness
 - All users should get access to CPU
 - Amount of CPU should be roughly even?
- Issue
 - Short-term vs. long-term fairness

21

Goals and Assumptions

- Goals (Performance metrics)
 - Minimize turnaround time: avg time to complete a job
 - Maximize throughput: operations (jobs) per second
 - Minimize overhead of context switches: large quanta
 - Efficient utilization (CPU, memory, disk etc)
 - Short response time: type on a keyboard
 - Small quanta
 - Fairness (fair, no starvation, no deadlock)
- Goals often conflict
 - Response time vs. throughput
 - fairness vs. avg turnaround time?
- Assumptions
 - One process/program per user
 - Programs are independent

22

Scheduling policies

- Is there an optimal scheduling policy?
- Even if we narrow down to one goal?
- But we don't know about future
 - Offline vs. online

23

Queuing Theory:

the mathematical study of waiting lines, or queues.

- An entire discipline to itself
- Mathematically oriented
- Some neat results
- Assumptions may be too restrictive to be able to model real-world situations exactly
 - E.g. assume infinite number of customers, infinite queue capacity, or no bounds on inter-arrival or service times
- Systems have grown more complex these days
- (Workload-driven) Simulation used instead now

24

CPU Scheduling policies

- FIFO
- Round Robin
- SJCF
- SRTCF

(break)


25

(Non-Preemptive scheduling) FIFO (FCFS) Policy

- What does it mean?
 - Run to completion (old days)
 - Run until blocked or yield
- Advantages
 - Simple
- Disadvantage?

26

Round Robin



- Each runs a time slice or **quantum**
- How do you choose time slice?
 - Overhead vs. response time
 - Overhead is typically about 1% or less
 - Quantum typically between 10 ~ 100 millisec

27

Is Fairness Always Good?

- Assume 10 jobs, each takes 100 seconds
- Assume no other overhead
- Total CPU time? 1000 seconds, always
- Implications?
 - Last job always finishes at 1000 seconds
 - So what's the point of scheduling?

28

Adding I/O Into the Mix

- Resource utilization example
 - A and B each uses 100% CPU
 - C loops forever (1ms CPU and 10ms disk)
 - Time slice 100ms: nearly 5% of disk utilization with Round Robin
 - Time slice 1ms: nearly 90% of disk utilization with Round Robin and nearly 100% of CPU utilization
- What do we learn from this example?
 - *Small time slice can improve utilization / fairness to I/O jobs*



29

Adding I/O Into the Mix


- Can we improve fairness for I/O bound processes without using tiny time slice?



30

Virtual Round Robin

- To improve fairness for I/O bound processes
- Aux queue is FIFO
- I/O bound processes go to aux queue (instead of ready queue) to get scheduled
- Aux queue has preference over ready queue



31

STCF vs. SRTCF

- Shortest time to completion first (shortest job first)
 - Non-preemptive
- Shortest *remaining* time to completion first
 - Preemptive
- When (what job arrival scenario) are they the same?
- Advantage
 - Minimal average turnaround time
- Disadvantage
 - Can cause starvation; difficult to know the future



32

Deep thinking



- Can you do better than Shortest job first in terms of average turnaround time?
- Prove it!

33

Read Chapter 5



34