

## CPU Scheduling (cont)

ECE595

Sep 18

Y. Charlie Hu

1

## Midterm

- Oct 4, 2018: 7-9pm
- Location: TBD

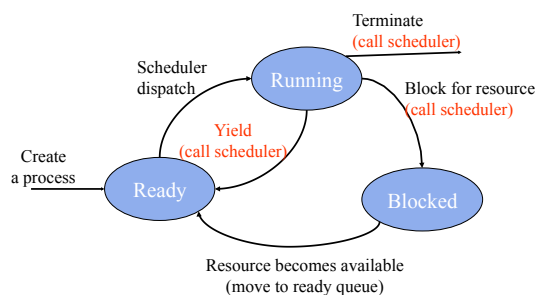
2

## Timesharing Systems

- **Timesharing** systems support interactive use
  - each user feels he/she has the entire machine
    - CPU, memory, I/O devices
- CPU - How?
  - optimize response time
  - based on time-slicing

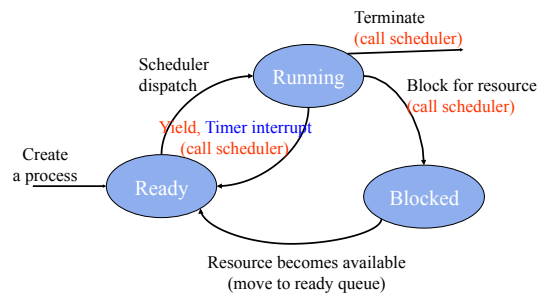
3

## Process State Transition of Non-Preemptive Scheduling



4

## Preemptive Scheduling



5

## Review on CPU scheduling

- Mechanism is easy, policy is hard
  - Jobs have diverse characteristics
  - 4 performance metrics
  - Don't know about future
- Hard to analyze even when narrowing down metric/job nature
  - FIFO vs. Round Robin in average turnaround time
    - 10 procs, 100 sec each, who wins?
    - 10 procs, 100 sec for first, 10 sec other 9, who wins?
  - Adding I/O to the diversity
- STCF vs. SRTCF

6

## STCF vs. SRTCF

- Shortest time to completion first (shortest job first)
  - Non-preemptive
- Shortest *remaining* time to completion first
  - Preemptive
- When (what job arrival scenario) are they the same?
- Advantage
  - Minimal average turnaround time
- Disadvantage
  - Can cause starvation; difficult to know the future

8

## Deep thinking

- Can you do better than Shortest job first in terms of average turnaround time?
- Prove it!

9

## Real challenge in using STCF



- Used in long-term scheduling batch system
  - e.g., explicit-length job queues (for repeated jobs)
- What about short-term scheduling (e.g. time-sharing system)?
  - Try prediction for repeated jobs
  - Exponential average:
$$\text{Pred}(n+1) = \alpha \cdot t_n + (1-\alpha) \cdot \text{Pred}(n)$$
- What about starvation?
  - Aging

10

## Observations so far



- Need to accommodate interactive jobs
  - Need some kind of RR
- Diversity in jobs – job length, I/O mix
  - RR also appears to help
- SJF also has virtue
  - Reduce avg. turnaround time
- Can we accommodate all?

11

## Scheduling policies



FIFO	Response time
	Throughput
RR	
	Avg. turnaround time
SJF	Fairness

12

## Break



True or false?

- “A CPU scheduling algorithm that minimizes avg. turnaround time cannot lead to starvation.”
- “Among all CPU scheduling algorithms, Round Robin always gives the worse average turnaround time.”

13

## Priority Scheduling



- To accommodate the spirits of SJF/RR/FIFO
- The method
  - Assign each process a *priority*
  - Run the process with highest priority in ready queue first
    - Use FIFO for processes with equal priority
  - Adjust priority dynamically
    - To deal with *all* issues: e.g. aging, I/O wait raises priority
- Advantage
  - Flexibility: Not all processes are “born” equal
- Challenge?

14

## Priority Scheduling (cont)



- Who sets the priorities?
  - Internally by OS
    - I/O to computation ratio (can be dynamic)
    - Memory requirement (can be dynamic)
    - Time constraints (e.g. real-time systems)
  - Externally by users/sysadm
    - Importance
    - Funds paid for
    - e.g. nice
- How -- Dynamically adjustment is tricky

15

## Approach 1: Multiple Queue Scheduling



- Motivation: processes may be of different nature and can be easily classified
  - e.g. foreground jobs vs. background jobs
- The method:
  - Processes permanently assigned to one queue, based on processes priority / type
  - Each queue has its own scheduling algo.
    - e.g. RR for foreground queue, FCFS for background queue
  - Need a scheduling among the queues
    - e.g. fixed priority preemptive scheduling (high-pri queue trumps other)
    - e.g. time-slice between queues

16

## Pros/Cons of Multiple Queue Scheduling



- Pros:
  - Low scheduling overhead
    - Jobs do not move across queues
- Cons:
  - Processes permanently assigned to one queue – not flexible
    - Program behavior may change
    - E.g. can switch between I/O bound and CPU bound
      - Need some learning/adaptation at runtime
  - Starvation cannot be easily handled
    - Need some learning/adaptation at runtime

17

## Approach 2: Multilevel Feedback Queue

	Priority	Time slices
—	0	1
—	1	2
—	2	4
—	3	8

- Jobs start at high priority queue
- Feedback
  - If timeout expires, drop one level
  - If waiting too long, push up one level (aging)
- Leave I/O bound and interactive processes in higher-priority queue

18

## UNIX System V

- One of the first commercial versions of the Unix
- Originally developed by AT&T
- First released in 1983
- 4 major versions, termed R1 R2, R3, and R4
- SVR4, commercially most successful version, the result of Unix System Unification (with collaborations of major UNIX vendors)
  - IBM's AIX is based on SVR3
  - Sun's Solaris and HP's HP-UX are based on SVR4
- 80s-early 90s, System V and BSD were the two major "flavors" of UNIX
  - Large multi-user systems vs. Desktop workstations

19

## Traditional Unix Scheduling (SVR3, 4.3 BSD)

Multilevel Feedback Queue with 1 sec preemption

- "1 sec" preemption
  - Preempt if a process doesn't block or complete within 1 sec
- Priority is recomputed every sec (low # is higher)
 
$$P_i = \text{base} + \text{CPU}_{i-1} / 2 + \text{nice}, \text{ where } \text{CPU}_i = (\text{U}_i + \text{CPU}_{i-1}) / 2$$

base is the base priority of the process  
 $\text{U}_i$  is process utilization in interval  $i$ , nice in  $[-20, 20]$
- Base priorities
  - Virtual memory Swapper
  - Block I/O device control
  - Character I/O device control
  - User processes

20

## Scheduling Algorithms in OSes

Operating System	Preemption	Algorithm
<a href="#">Windows 3.1x</a>	None	<a href="#">Cooperative Scheduler</a>
<a href="#">Windows 95, 98, Me</a>	Half	Preemptive for 32-bit processes, <a href="#">Cooperative Scheduler</a> for 16-bit processes
<a href="#">Windows NT</a> (2000, XP, Vista, 7, and Server)	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">Mac OS</a> pre-9	None	<a href="#">Cooperative Scheduler</a>
<a href="#">Mac OS 9</a>	Some	Preemptive for MP tasks, <a href="#">Cooperative Scheduler</a> for processes and threads
<a href="#">Mac OS X</a>	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">Linux</a> pre-2.6	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">Linux</a> 2.6-2.6.23	Yes	<a href="#">O(1) scheduler</a>
<a href="#">Linux</a> post-2.6.23	Yes	<a href="#">Completely Fair Scheduler</a>
<a href="#">Solaris</a>	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">NetBSD</a>	Yes	<a href="#">Multilevel feedback queue</a>
<a href="#">FreeBSD</a>	Yes	<a href="#">Multilevel feedback queue</a>

22

## Cooperative Scheduling

- Early multitasking systems used applications that **voluntarily** ceded time to one another.
- This approach, which was eventually supported by many OSes, is known today as cooperative multitasking.
- Cooperative multitasking was once the scheduling scheme employed by [Microsoft Windows](#) (prior to [Windows 95](#) and [Windows NT](#)) and [Mac OS](#) (prior to [Mac OS X](#)) to enable multiple applications to be run simultaneously
- Now rarely used in larger systems

23

## Example Details in Chapter 5

- Solaris 2
- Windows XP
- Linux
- All are variations of Multilevel Feedback Queue


24

## Mechanism vs policy: Lottery Scheduling

- Motivation
  - Priority scheduling is implemented by adjusting priorities
  - Adjusting priority is a bit ad hoc (e.g. what rate?)
- Lottery method [OSDI 94]
  - Give each job a number of tickets
  - High priority jobs get more tickets
  - Randomly pick a winning ticket
  - To avoid starvation, give each job at least one ticket


25

## Multiprocessor and Cluster



Multiprocessor architecture

- L2 cache coherence
- A single “image” OS



Cluster/Multicomputer

- Distributed memory
- An OS on each box

27

## Multiprocessor/Cluster Scheduling



- New design issue: process/thread inter-dependence
  - Threads of the same process may synchronize
  - Processes of the same job may send/recv messages

28

## Multiprocessor/Cluster Scheduling: Example Approach



- Gang scheduling (coscheduling)
  - Threads of same process will run together on multiprocessor
  - Processes of same application run together on cluster

29

## Reading assignment



- OSC, Chapter 5

30