

Sharing Main Memory, Segmentation

ECE595
Oct 2

Y. Charlie Hu



1


Midterm

- Time: Oct 4, Thursday, 7-9pm
- Location: EE 222



3

Connecting the dots



4

The big picture

- a.out needs address space for
 - text seg, data seg
- A running process needs phy. memory for
 - text seg, data seg, heap, stack
- But no way of knowing where in phy mem at
 - Programming time, compile time, linking time
- **Best way out?**
 - Make agreement to divide responsibility
 - Assume address starts at 0 at prog/compile/link time
 - OS needs to work hard at loading/running time




5

Big picture (cont)

- OS deals with physical memory
 - Loading
 - Sharing physical memory between processes
 - Dynamic memory allocation

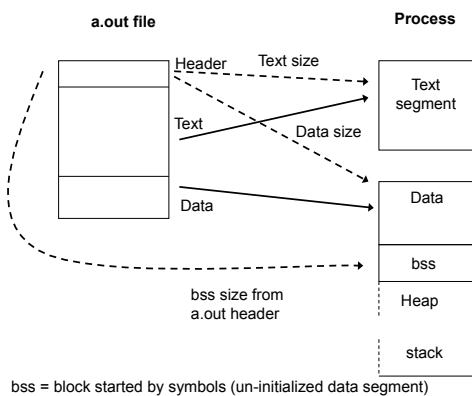
6

Connecting the dots



7

Loading




8

Dynamic memory allocation during program execution

- Stack: for procedure calls
- Heap: for malloc()
- Both dynamically growing/shrinking
- Assumption for now:
 - Heap and stack are fixed size
 - OS has to worry about loading 4 segments per process:
 - Text
 - Data
 - Heap
 - stack

9

1. Simple uniprogramming: Single segment (code, data, stack heap) per process



10

Simple uniprogramming: Single segment per process

- Highest memory holds OS
- Process is allocated memory starting at 0, up to the OS area
- When loading a process, just bring it in at 0
 - virtual address == physical address!
- Examples:
 - early batch monitor which ran only one job at a time
 - if the job wrecks the OS, reboot OS
 - 1st generation PCs operated in a similar fashion
- Pros / Cons?

11

Multiprogramming

- Want to let several processes coexist in main memory

12

Issues in sharing main memory

- **Transparency:**
 - Processes should not know memory is shared
 - Run regardless of the number/locations of processes
- **Safety:**
 - Processes mustn't be able to corrupt each other
- **Efficiency:**
 - Both CPU and memory utilization shouldn't be degraded badly by sharing

13


2. Simple multiprogramming

With **static software memory relocation**, no protection, 1 segment per process:

- Highest memory holds OS
- Processes allocated memory starting at 0, up to the OS area
- When a process is loaded, **relocate** it so that it can run in its allocated memory area
 - How? (use symbol table and relocation info)
- Analogy to linking?


14

Simple multiprogramming: Single segment per process, static relocation



15

Simple multiprogramming: Single segment per process, static relocation



16

Simple multiprogramming: Single segment per process, static relocation

- 4 drawbacks
 1. No protection
 2. Low utilization -- Cannot relocate dynamically
 - Binary is fixed (after loading)
 - Cannot do anything about holes
 3. No sharing -- Single segment per process
 - Cannot share part of process address space (e.g. text)
 4. Entire address space needs to fit in mem
 - Need to swap whole, very expensive!

17

What else can we do?

- Already tried
 - Compile time / linking time
 - Loading time
- Let us try execution time!

18


3. Dynamic memory relocation

- Instead of changing the address of a program before it's loaded, change the address dynamically *during every reference*
 - Under dynamic relocation, each program-generated address (called a *logical address* or *virtual address*) is translated in hardware to a *physical* or *real address*

Can this be done in software?

19

Translation overview




- Actual translation is in hardware (MMU)
- Controlled in software
- CPU view
 - what program sees, virtual addresses
- Memory view
 - physical memory addresses

20

break

21


3.1 Base and bound



- Built in Cray-1 (1976)
- A program can only access physical memory in [base, base+bound]
- On a context switch: save/restore base, bound registers
- Pros:
 - simple, fast, cheap
 - Can relocate segment

22


3.1 Base and bound



- The essence:
 - A level of indirection
 - $\text{Phy. Addr} = \text{Vir. Addr} + \text{base}$

23

Base and bound



- Cons:
 - Only one segment
 - How can two processes share code while keeping private data areas (shared editors)?
 - Can it be done safely with a single-segment scheme?


24

What have we achieved?

- 4 drawbacks
 1. No protection
 2. Low utilization -- Cannot relocate dynamically
 - Cannot do anything about holes
 3. No sharing -- Single segment per process
 - Cannot share part of process address space (e.g. text)
 4. Entire address space needs to fit in mem
 - Need to swap whole, very expensive!

25

3.2 Multiple Segments



- Have a table of (seg, size)
- Further protection: each entry has (nil, read, write, exec)
- On a context switch: save/restore the table (or a pointer to the table) in kernel memory

26

How does this allow 2 processes to share code segment?

27

Segmentation example

text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

01C2: jmp 01F4

01E0: call 0320

01F4: X:

0320: bar:

Data segment [0x1000, 0x16A0]

15DC: _Y:

2-bit segment number, 12-bit offset

Segment	Base	Bounds	RW
0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

28

Segmentation example

text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

01C2: jmp 01F4

01E0: call 0320

01F4: X:

0320: bar:

Data segment [0x1000, 0x16A0]

15DC: _Y:

2-bit segment number, 12-bit offset

Segment	Base	Bounds	RW
0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

29

→ Where is 01F4 in physical memory?

Segmentation example



text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x16A0]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

→ Where is 15DC in physical memory?

30

Segmentation example



text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x16A0]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0	4000	4B0	10
1	0	6A0	11
2	3000	FFF	11
3	--	--	00

→ Suppose SP is initially 265C. Where is it in physical mem?³¹

Segmentation example



text segment [0x0000, 0x04B0]

foo: bar procedure

019A: LD R1, 15DC

0320: bar:

01C2: jmp 01F4

01E0: call 0320

Data segment [0x1000, 0x1190]

01F4: X:

15DC: _Y:

2-bit segment number, 12-bit offset

Segment Base Bounds RW

0	4000	4B0	10
1	0	190	11
2	3000	FFF	11
3	--	--	00

→ which portions of the virtual and physical address spaces are used by this process?

32

Pros/cons of segmentation



Pros:

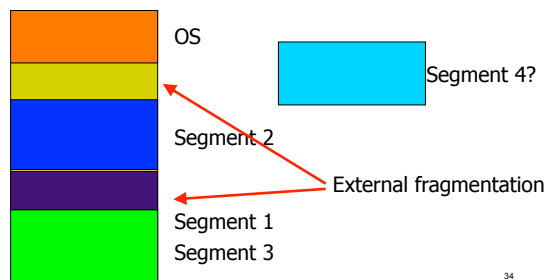
- Process can be split among several segments
 - Allows sharing
- Segments can be assigned, moved, or swapped independently

Cons:

- External fragmentation:** many holes in physical memory
 - Also happens in base and bound scheme

33

Simple multiprogramming: Single segment per process, static relocation



34

What fundamentally causes external fragmentation?

1. Segments of many different sizes
2. Each has to be allocated contiguously

35

Dynamic memory allocation problem

- Problem: External fragmentation caused by holes too small
- How much can a smart allocator help?
 - The allocator maintains a free list of holes
 - Allocation algorithms differ in how to allocate from the free list

36

Dynamic allocation algorithms

- **Best fit:** allocate the smallest chunk big enough
- **First fit:** allocate the first chunk big enough
 - Rotating first fit
- Is best fit necessarily better than first fit?
 - Example: 2 free blocks of size 20 and 15
 - If allocation requests are 10 then 20, which one wins?
 - If requests are 8, 12, then 12, which one wins?

37

Dynamic allocation algorithms



- Analysis shows
 - First fit tends to leave average-size holes
 - Best fit tends to leave some very large holes, very small holes
- Knuth claims that if storage is close to running out, it will run out regardless of which scheme is used
 - Pick the easiest or most efficient (e.g. first fit)

38

Segmentation: OS implementation



- Keep segment table in PCB
- When creating process, allocate space for segments, fill in PCB bases/bounds
- When process dies, return physical space used by segments to free pool
- Context switch?
 - Saves old segment table / Loads new segment table to MMU

39

[lec3] Kernel data structure: Process Control Block (Process Table)



- Process management info
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- Memory management info
 - Segments, page table, stats, etc
- I/O and file management
 - Communication ports, directories, file descriptors, etc.

40

Managing segments



To enlarge a segment:

- See if space above segment is free. If so, just update the bound and use that space
- Or, move this segment to disk and bring it back into a larger hole (or maybe just copy it to a large hole)

41

Managing segments (cont)

- When there is no space to allocate a new segment:
 - Compact memory – how?

42

Summary: Evolution of Memory Management (so far)

Scheme	How	Pros	Cons
Simple uniprogramming	1 segment loaded to starting address 0	Simple	1 process 1 segment No protection
Simple multiprogramming	1 segment relocated at loading time	Simple, Multiple processes	1 segment/process No protection External frag.
Base & Bound	Dynamic mem relocation at runtime	Simple hardware, Multiple processes Protection	1 segment/process, External frag.
Multiple segments	Dynamic mem relocation at runtime	More hardware, Protection, multi segs/process	External frag.

43

Reading assignment

- Read Chapter 8

44