

Demand Paging


ECE595

Oct 18

Y. Charlie Hu

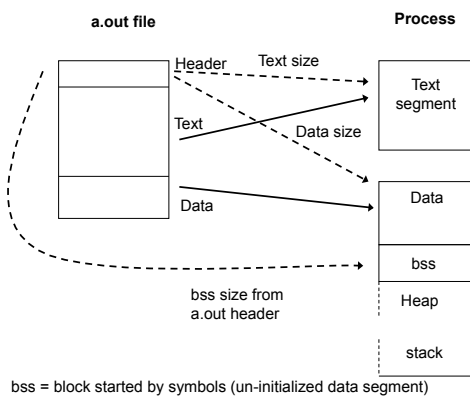
1

[lec13] The big picture – connecting the dots



2


[lec13] Loading



3

Review: sharing main memory

- Simple multiprogramming



4 drawbacks?

4

[lec14] Sharing main memory

- Simple multiprogramming – 4 drawbacks
 - Lack of protection
 - Cannot relocate dynamically
 - Single segment per process
 - Entire address space needs to fit in mem

Dynamic
Memory
Relocation


5

[lec15] 3. Dynamic memory relocation

- Instead of changing the address of a program before it's loaded, change the address dynamically *during every reference*
 - Under dynamic relocation, each program-generated address (called a *logical address* or *virtual address*) is translated in hardware to a *physical* or *real address* at runtime

6

[lec15] Translation overview



- Actual translation is in hardware (MMU)
- Controlled in software
- CPU view
 - what program sees, virtual memory
- Memory view
 - physical memory

What is the essence of going through MMU?

7

Sharing main memory

- Simple multiprogramming – 4 drawbacks
 - Lack of protection
 - Cannot relocate dynamically
 - Base-and-bound
 - Single segment per process
 - Paging, segmentation, etc.
 - Entire address space needs to fit in mem

Dynamic
Memory
Relocation

8

Sharing main memory

- Simple multiprogramming – 4 drawbacks
 - Lack of protection
 - Cannot relocate dynamically
 - dynamic memory relocation: base& bound
 - Single segment per process
 - dynamic memory relocation: segmentation, paging
- Entire address space needs to fit in mem
 - More need for swapping
 - Need to swap whole, very expensive!

9

The last drawback

- So far we've separated the process's view of memory from the OS's view using a mapping mechanism
 - Each sees a different organization
 - Allows OS to shuffle processes around
 - Simplifies memory sharing
 - *What is the essence of the mechanism that enables this?*
- But, a user process had to be completely loaded into memory before it could run
 - Wasteful since a process only needs a small amount of its total memory at any time (*reference locality!*)

10

[lec1] What is an OS?

Extended (abstract) machine (answer 2)

- Much more ideal environment than the hardware
 - Ease to use
 - Fair (well-behaved)
 - Portable (back-compatible)
 - Reliable
 - Safe
- Illusion of infinite, private resources
 - Single processor → many separate processors
 - *Single memory → many separate, larger memories*

11

Virtual Memory

- Definition: *Virtual memory* permits a process to run with only some of its virtual address space loaded into physical memory
- Key idea: Virtual address space translated to either
 - Physical memory (small, fast) or
 - Disk (backing store), large but slow
- [Deep thinking] What made above possible?
- Objective:
 - To produce the illusion of memory as big as necessary

12

Virtual Memory

- “To produce the illusion of memory as big as necessary”
 - Without suffering a huge slowdown of execution
 - Why makes this possible?
 - *Principle of locality*
 - Knuth’s estimate of 90% of the time in 10% of the code
 - There is also significant locality in data references

13

Virtual Memory Implementation


- Virtual memory is typically implemented via *demand paging*
- *demand paging*: paging with swapping, e.g., physical pages are swapped in and out of memory

14

Demand Paging (paging with swapping)

- If not all of a program is loaded when running, what happens when referencing a byte not loaded yet?

- How to *detect* this?
 - In software?



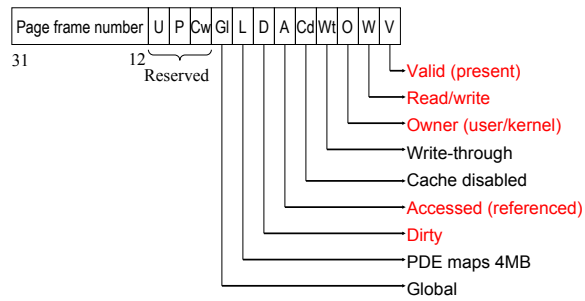
15

Demand Paging (paging with swapping)

- If not all of a program is loaded when running, what happens when referencing a byte not loaded yet?
- Hardware/software cooperate to make things work
 - Extend PTEs with an extra bit “*present*”
 - Any page not in main memory right now has the “present” bit cleared in its PTE
 - If “present” isn’t set, a reference to the page results in a trap by the paging hardware, called *page fault*
 - What needs to happen when page fault occurs?


16

x86 Page Table Entry



17

Page Fault Handling in Demand Paging



18

Page fault handling (cont)

- On a page fault
 - Find an unused phy. page or a used phy. page (how?)
 - If the phy. page is used
 - If it has been modified, write it to disk
 - Invalidate its current PTE and TLB entry (how?)
 - Load the new page from disk
 - Update the faulting PTE and its TLB entry
 - Restart the faulting instruction
- Supporting data structure
 - For speed: A list of unused physical pages (more later)
 - Data structure to map a phy. page to its pid and virtual address
 - Sounds familiar?

19

Page Fault Handling: Subtlety

- Page fault may have occurred in the middle of an instruction!
 - Suppose the instruction is restarted from the beginning
 - How is beginning located?
 - Side effects? MOVE (SP)+, R2
 - Require hardware support to keep track of side effects of instructions and undo them before restarting
- You can make it easy when designing ISA
 - RISC (load-store) architecture make this relatively easy – only load/store can generate page faults

20

Deep thinking: VM implementation



- Virtual memory is typically implemented via *demand paging*
- It can also be implemented via *demand segmentation*
 - Double drawbacks?

22

Deep thinking: Virtual Address vs. Virtual Memory



- Virtual address
 - Supported with dynamic memory relocation
 - Segmentation
 - Paging
- Virtual memory
 - Dynamic memory relocation + *swapping*
 - *Demand* paging
 - *Demand* segmentation

23