

TLB. Mac OS Memory Management

ECE595

Oct 23

Y. Charlie Hu

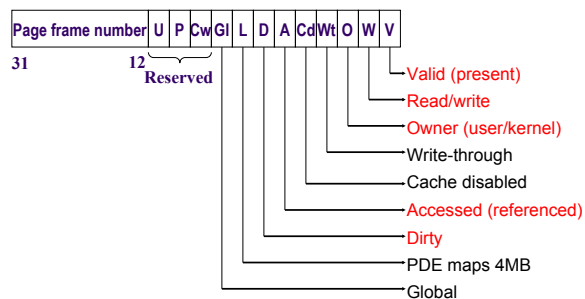
1

[lec15] What is happening before and after malloc()?

- Before malloc()?
- After malloc()?
- Upon first access?
- How to capture the first write to a virtual page?
 - e.g. want to trap into page fault handler
 - Use valid bit
 - In handler, check if vpage is malloced.
 - If not, segmentation fault
 - Else allocate physical page

2

x86 Page Table Entry



3

[lec16] Virtual Memory

- Definition: *Virtual memory* permits a process to run with only some of its virtual address space loaded into physical memory
- Key idea: Virtual address space translated to either
 - Physical memory (small, fast) or
 - Disk (backing store), large but slow
- [Deep thinking] What made above possible?
- Objective:
 - To produce the illusion of memory as big as necessary

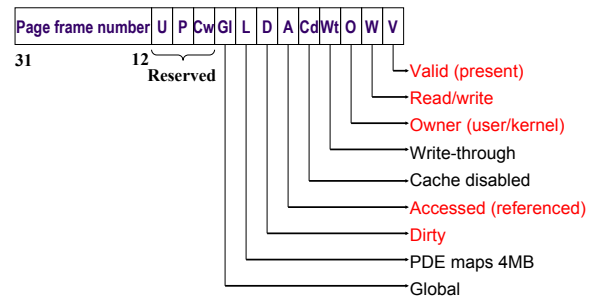
4

Demand Paging (paging with swapping)

- If not all of a program is loaded when running, what happens when referencing a byte not loaded yet?
- Hardware/software cooperate to make things work
 - Extend PTEs with an extra bit “**present**”
 - Any page not in main memory right now has the “present” bit cleared in its PTE
 - If “present” isn’t set, a reference to the page results in a trap by the paging hardware, called **page fault**
 - What needs to happen when page fault occurs?


5

x86 Page Table Entry



6

Page Fault Handling in Demand Paging



7


Today's topics

- TLB
- Mac OS Memory Management

8

[lec14] Paging implementation – how does it really work?

Virtual address



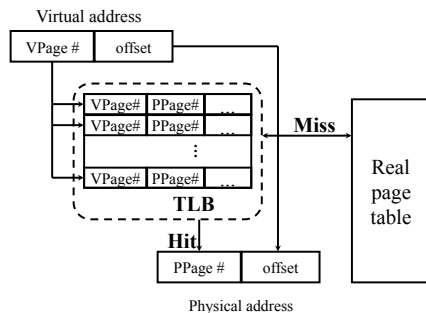
- Where to store page table?
- How to use MMU?
 - Even small page tables too large to load into MMU
 - Page tables kept in mem and MMU only has their base addresses
 - What does MMU have to do?
- Page size?
 - Small page -> big table
 - 32-bit with 4k pages
 - Large page -> small table but large internal fragmentation
 - e.g., data seg is 6 Kbytes

Performance problem with paging

- How many extra memory references to access page tables?
 - One-level page table?
 - Two-level page table (midterm problem)?
- Solution: *reference locality!*
 - In a short period of time, a process is likely accessing only a few pages
 - Instead of storing only page table starting address in MMU,

10


Translation Look-aside Buffer (TLB)



TLB often fully set-associative → least conflict misses
Expensive → typically 64 – 1024 entries

11

Bits in a TLB Entry



- Common (necessary) bits
 - Virtual page number: match with the virtual address
 - PTE
- Optional (useful) bits
 - ASIDs -- Address-space identifiers (process tags)

12

Miss handling: Hardware-controlled TLB



- On a TLB hit, MMU checks the valid bit
 - If valid, perform address translation
 - If invalid (e.g. page not in memory), MMU generates a page fault
 - OS performs fault handling
 - Restart the faulting instruction
- On a TLB miss
 - MMU parses page table and loads PTE into TLB
 - Needs to replace if TLB is full
 - PT layout is **fixed**
 - Same as hit ...

13

Miss handling: Software-controlled TLB



- On a TLB hit, MMU checks the valid bit
 - If valid, perform address translation
 - If invalid (e.g. page not in memory), MMU generates a page fault
 - If page not valid, OS performs page fault handling
 - Restart the faulting instruction
- On a TLB miss, HW raises exception, **traps to the OS**
 - OS parses page table and loads PTE into TLB
 - Needs to replace if TLB is full
 - PT layout is **flexible**
 - Same as in a hit...

14

Hardware vs. software controlled



- | | |
|---|--|
| <ul style="list-style-type: none"> • Hardware approach <ul style="list-style-type: none"> • Efficient – TLB misses handled by hardware • OS intervention is required only in case of page fault • Page structure prescribed by MMU hardware -- rigid | <ul style="list-style-type: none"> • Software approach <ul style="list-style-type: none"> • Less efficient -- TLB misses are handled by software • MMU hardware very simple, permitting larger, faster TLB • OS designer has complete flexibility in choice of MM data structure <ul style="list-style-type: none"> • e.g. 2-level page table, inverted page table) |
|---|--|

15

Deep thinking



- Without TLB, how MMU finds PTE is fixed
- With TLB, it can be flexible, e.g. software-controlled is possible
- What enables this?

16

More TLB Issues



- Which TLB entry should be replaced?
 - Random
 - LRU
- What happens when changing a page table entry (e.g. because of swapping, change read/write permission)?
 - change the entry in memory
 - flush (eg. invalidate) the TLB entry
 - INGLPG on x86

17

What happens to TLB in a process context switch?



- During a process context switch, cached translations can not be used by the next process
 - Invalidate all entries during a context switch
 - Lots of TLB misses afterwards
 - Tag each entry with an ASID
 - Add a HW register that contains the process id of the current executing process
 - TLB hits if an entry's process id matches that reg

18

Cache vs. TLB



- Similarities:
 - Both cache a part of the physical memory
- Differences:
 - Associatively
 - TLB is usually fully associative
 - Cache can be direct mapped
 - Consistency
 - TLB does not deal with consistency with memory
 - TLB can be controlled by software

19

More on consistency Issues




- Snoopy cache protocols can maintain consistency with DRAM, even in the presence of DMA
- **No hardware maintains consistency between DRAM and TLBs:**
 - OS needs to flush related TLBs whenever changing a page table entry in memory
- On multiprocessors, when you modify a page table entry, you need to do “**TLB shoot-down**” to flush all related TLB entries on all processors

20

Memory Hierarchy Revisited

What does this imply about L1 addresses?

Where do we hope requests get satisfied?




21

Memory Hierarchy Re-Revisited

What does this imply about L1 addresses?

Any speed benefits?
Any drawbacks?



22

Today's topics

- TLB
- Mac OS Memory Management

23

Instances of extra level of indirection

- Dynamic memory relocation
 - Base and bound
 - Segmentation
 - 1-level paging
 - 2-level paging
 - Segmentation plus paging
- The MAC example

24

Classic example of indirection: MAC OS memory management



- Macintosh had 128 KB (~1985)
- Ran one app at a time
- Memory constraint was high priority
- Solution?
 - willing to trade off running time
 - “All computer science problems can be solved with an extra level of indirection.”

25

Reading



- Chapter 8

26