

Page Replacement Algorithms

ECE595
Oct 25

Y. Charlie Hu



1

Virtual Memory review



- What is *Virtual Memory*?
- Virtual address vs. virtual memory?
- How is Virtual Memory typically implemented?
 - What is *demand paging*?
 - It can also be implemented via *demand segmentation*
 - Double drawbacks?
- In Virtual Memory, what will the virtual address space be translated to?
- What is a page fault?
- Page fault vs. TLB miss?
- Does a TLB hit imply no page fault?

2

[lec16] Virtual Memory



- Definition: *Virtual memory* permits a process to run with only some of its virtual address space loaded into physical memory
- Virtual address space translated to either
 - Physical memory (small, fast) or
 - Disk (backing store), large but slow
- Objective:
 - To produce the illusion of memory as big as necessary

3

[lec16] Virtual Memory



- “To produce the illusion of memory as big as necessary” (yet as fast as physical memory)
 - What makes this possible?

4

[Ice1] Separating Policy from Mechanism

A fundamental design principle in Computer Science

Mechanism – tool/implementation to achieve some effect

Policy – decisions on what effect should be achieved

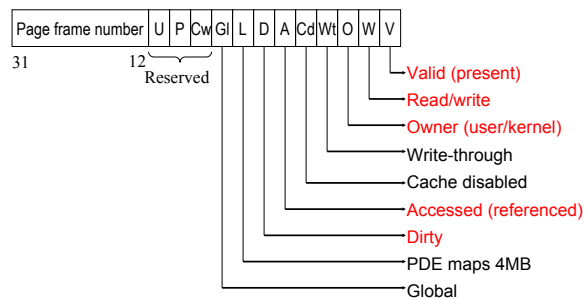
Examples:

- All users treated equally
- All program instances treated equally
- Preferred users treated better

Separation leads to flexibility

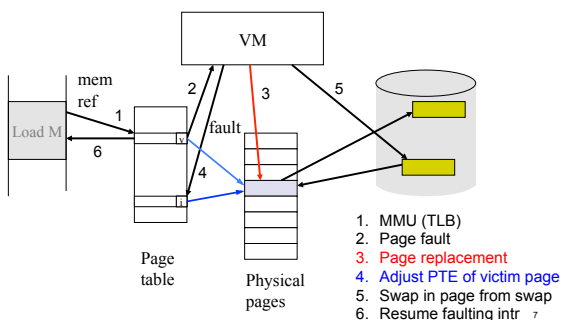
5

The Mechanism: x86 Page Table Entry



6

The Mechanism: Page Fault Handling in demand paging



7

Page out on critical path?

- Page in has to wait till page out is finished
 - Page fault handling time = proc. overhead + 2 * I/Os
- Solution: page buffering
 - System maintains a pool of free pages
 - When a page fault occurs, victim page chosen as before
 - But desired page paged into a free page right away before victim page paged out
 - Victim page written out when disk is idle

8

The Policy Issue: Page selection and replacement

- Once the hardware has provided basic capabilities for VM, the OS must make two kinds of decisions
 - Page selection:** *when* to bring pages into memory
 - Page replacement:** *which* page(s) should be thrown out
 - Try to kick out the least useful page

9

Page selection (when)

- Prepaging
 - Bring a page into memory before it is referenced
 - Hard to do without a "prophet"
- Request paging
 - Let user say which pages are needed when
 - Users don't always know best
 - And aren't always impartial
- Demand paging
 - Start up process with no pages loaded
 - Load a page when a page fault occurs, i.e., wait till it **MUST** be in memory
 - Almost all paging systems are demand paging

10


Page replacement problem

- Definition
 - Expect to run with all physical pages in use
 - Every "page-in" requires an eviction to swap space
 - How to select the victim page?
- Performance goal:
 - Give a sequence of memory accesses, minimize the # of page faults
 - given a sequence of virtual page references, minimize the # of page faults
- Intuition: kick out the page that is least useful
- Challenge: how do you determine "least useful"?
 - Even if you know future?

11

What makes finding the least useful page hard?

- Don't know future
 - Using past to predict future is pretty accurate
- Getting perfect reference stream is hard (why?)
- Can we get imperfect information without high cost?



12

Optimal or MIN


- Algorithm:
 - Replace the page that won't be used for *the longest time*
- Pros
 - Minimal page faults (can you prove it?)
 - This is an **off-line** algorithm for performance analysis
- Cons
 - No **on-line** implementation
- Also called **Belady's Algorithm**
- What was the CPU scheduling algorithm of similar nature?
 - Why is the proof of MIN so much harder?

13

What can we do without extra hardware support?

14

First-In-First-Out (FIFO)



- Algorithm
 - At page fault, throw out the oldest page
- Pros?
- Cons?

15

Predicting future based on past

- “Principle of locality”
 - Recency:
 - Page **recently** used are likely to be used again in the near future
 - Frequency:
 - Pages **frequently** used (recently) are likely to be used frequently again in the near future
- Is this temporal or spatial locality?
- Why not spatial locality?

16

How to record locality in the past?

- Software solution?



17

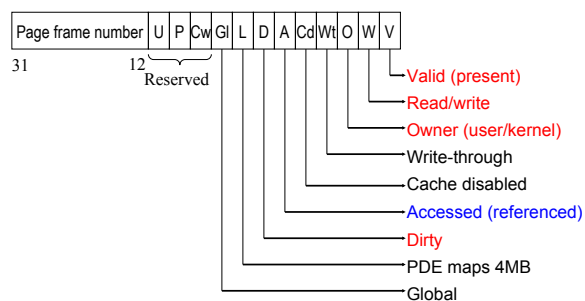
Exploiting locality needs some hardware support

- **Reference bit**
 - A hardware bit that is set whenever the page is referenced (read or written)




18

x86 Page Table Entry



19

FIFO with 2nd Chance




- **Algorithm**
 - Check the reference-bit of the oldest page
 - If it is 0, then replace it
 - If it is 1, clear the referent-bit, put it to the end of the list, and continue searching
- **Pros**
 - Potentially fast
 - Frequency → do not replace a “heavily” used page
- **Cons**
 - The worst case may take a long time



20


Clock: a simple FIFO with 2nd chance



- FIFO clock algorithm
 - Hand points to the oldest page
 - On a page fault, follow the hand to inspect pages
- Second chance
 - If the reference bit is 1, set it to 0 and advance the hand
 - If the reference bit is 0, use it for replacement
- What is the difference between Clock and the previous one?
 - Mechanism vs. policy?

21

Clock: a simple FIFO with 2nd chance



- What happens if all reference bits are 1?
- What does it suggest if observing clock hand is sweeping very fast?
- What does it suggest if clock hand is sweeping very slow?

22

Can we refine the Clock Algorithm?

- Key observation: it is cheaper to pick a “clean” page over a “dirty” page
 - Clean page does not need to be swapped to disk
- Challenge:
 - How to get this info?

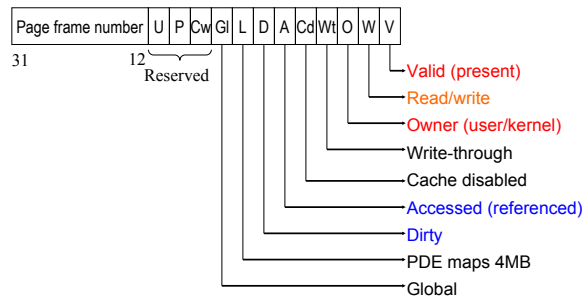
23

Refinement by adding extra hardware support

- **Reference bit**
 - A hardware bit that is set whenever the page is referenced (read or written)
- **Modified bit (dirty bit)**
 - A hardware bit that is set whenever the page is written into

24

x86 Page Table Entry



25

Enhanced FIFO with 2nd-Chance Algorithm (used in Macintosh VM)

- Same as the basic FIFO with 2nd chance, except that it considers both (reference bit, modified bit)
 - (0,0): neither recently used nor modified (good)
 - (0,1): not recently used but dirty (not as good)
 - (1,0): recently used but clean (not good)
 - (1,1): recently used and dirty (bad)
- When giving second chance, only clear reference bit
- Pros
 - Avoid write back
- Cons
 - More complicated, worse case scans multiple rounds

26

Enhanced FIFO with 2nd-Chance Algorithm – implementation

- On page fault, follow hand to inspect pages:
 - Round 1:
 - If bits are (0,0), take it
 - if bits are (0,1), record 1st instance
 - Clear ref bit if (0,1) not found yet
 - At end of round 1, if (0,1) was found, take it
 - If round 1 does not succeed, try 1 more round

27

FIFO revisited

- Book-keeping when pages are loaded in or kicked out
- No idea or vague idea (2 bits) what happened in between two 2 faults
- Ideally, what more info would we like?


28

Least Recently Used (LRU)

- Algorithm
 - Replace page that hasn't been used for the longest time
- Question
 - What hardware mechanisms are required to implement *exact* LRU?

29

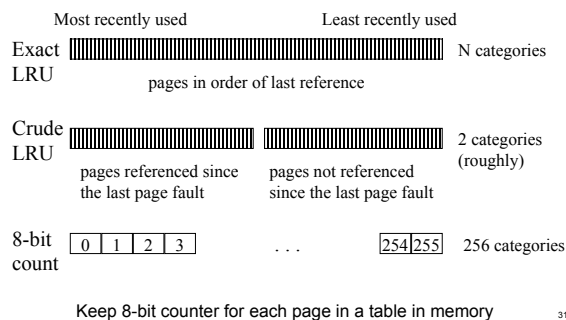
(exact) LRU implementation



- ideally
 - Extend PTE with a time-of-use field
 - Add to CPU a clock, incremented & copied into the PTE for *every* mem reference
 - So we have the time of the last reference to each page
 - At TLB eviction?
 - At finding the LRU page?

30

Approximate LRU



31

Approximate LRU

Interval 1	Interval 2	Interval 3	Interval 4	Interval 5
00000000	00000000	10000000	01000000	10100000
00000000	10000000	01000000	10100000	01010000
10000000	11000000	11100000	01110000	00111000
00000000	00000000	00000000	10000000	01000000

- Algorithm
 - At regular interval, OS shifts reference bits (in PTE) into counters (and clear reference bits)
 - Replacement: Pick the page with the "smallest counter"
- How many bits are enough?
 - In practice 8 bits are quite good
- Pros: Require one reference bit, small counter/page
- Cons: Require looking at many counters (or sorting)

32

Page replacement algorithms: Summary



- Optimal
- FIFO
- Random
- FIFO with 2nd chance
- Clock: a simple FIFO with 2nd chance
- Enhanced FIFO with 2nd chance
- Approximate LRU

33

Reading



- Chapter 9

34