


Introduction to File Systems: FS Structure, Disk Layout

ECE 595
Nov 1

Y. Charlie Hu




Disk



- Platter / Head / Tracks / Sectors / Cylinders
- $256 \text{ heads} * 17849 \text{ tracks} * 63 \text{ sectors} * 512 \text{ bytes} = 140 \text{ GB disk}$

Moving-head Disk Mechanism



Overview of Disk (non-volatile storage)

- Magnetic disks provide bulk of secondary storage of modern computers
 - Drives rotate at 60 to 200 times per second
 - **Transfer rate** is rate at which data flow between drive and computer
 - **Positioning time (random-access time)** is time to move disk arm to desired cylinder (**seek time**) and time for desired sector to rotate under the disk head (**rotational latency**)
 - **Head crash**: disk head making contact w/ disk surface, not repairable
- Disks can be removable (hot swap)
- Drive attached to computer via **I/O bus**
 - Buses vary, including **EIDE, ATA, SATA, USB, SCSI**
 - **Host controller** in computer uses bus to talk to **disk controller** built into drive or storage array



Disk formatting

- *Low-level formatting, or physical formatting* — Dividing a disk into sectors (logical blocks) that the disk controller can read and write
 - 512 bytes, 1024 bytes
- Handling bad blocks (extra-level of indirection!)
 - IDE drives: “manually” scan the disk and mask out bad blocks in file system management
 - SCSI drives: disk controller maintains and masks off bad blocks
 - *Sector sparing*
 - Low-level formatting sets aside spare sectors
 - Controller can replace bad blocks logically with a spare one

Disk Accesses


- Modern disk drives are addressed as large 1D arrays of **logical blocks**, where logical blocks are smallest transfer units
 - The size of logical blocks is usually 512 bytes
 - The logical blocks have high reliability (e.g. SCSI)
 - File system often has its own unit, called **file system block**, that is a multiple of disk blocks
 - e.g. 4K bytes

6

Disk



“My lab3 solution is in
~ece595/lab3/process.c.”



“I will save my lab4
solution on platter 5,
track 8739, sector 3,
offset 10 bytes.”

- Platter / Head / Tracks / Sectors / Cylinders
- $256 \text{ heads} * 17849 \text{ tracks} * 63 \text{ sectors} * 512 \text{ bytes}$
= 140 GB disk

File System

- The most visible aspect of an operating system
 - Files
 - Directories
 - Protections
 - Persistence
 - Transparent remote access

8

What is a File?



- File: a named collection of bytes stored on disk
 - From OS's standpoint
 - A file consists of a bunch of blocks stored on disk
 - From programmer's view
 - A collection of data records
 - File system performs the magic / translation
 - Pack bytes into disk blocks on writing
 - Unpack them again on reading

9

[lec1] Why Files?



- Physical reality
 - Block oriented
 - Physical sector numbers
 - No protection among users of the system
 - Data might be corrupted if machine crashes
- File system abstraction
 - Byte oriented
 - Named files
 - Users protected from each other
 - Robust to machine failures

10

File Types



- ASCII – plain text
 - inbox
- A Unix executable file (a.out)
 - header: magic number, sizes, entry point, flags
 - text (code)
 - data
- Devices
 - keyboard
 - terminal

11

Common Addressing Patterns




- *Sequential*: information is processed in order, one piece after another
 - Example?
- *Random (direct)*: can access any block in the file directly without passing through its predecessors
 - E.g. databases
- *Content based*: search for blocks with particular values
 - E.g. hash table, dictionary

12

So What Makes File Systems Hard?


- Files grow and shrink
 - Little *a priori* knowledge
 - 6~8 orders of magnitude in file sizes
- Overcoming disk performance behavior
 - Highly nonuniform access
 - Desire for efficiency
- Coping with failure



13

File System Components


- Naming/Access
 - User gives file name, not track or sector number, to locate data
- Disk management
 - Arrange collection of disk blocks into files
- Security
 - Keep information secure
- Reliability/durability
 - When system crashes, lose stuff in memory, but want files to be durable



14

Roadmap (The Journey)

- Functionality** (naming/access)
 - Data structures
 - Files/directories
 - File operations
- Performance** (disk management)
 - Disk layout (I/O performance)
 - Disk scheduling (I/O performance)



15

Definitions

- File descriptor** (fd) – an integer used to represent a file – easier than using names
- Metadata** – bookkeeping data that describes the file or info about it; not the actual content
 - inode** – “index node”, file metadata on Unix
 - inode design discussed in next lecture
- Open file table** – in-memory, system-wide list of file metadata in use

16


File System API

- OS provides the file system abstraction
- How do application processes access the file system?

17

Opening a File: `fd = open("file1")`

- File name lookup and authentication
- Create an entry in the open file table (system wide) if it is not in
- Copy the file metadata to in-memory data structure, if it is not in
- Create an entry in PCB
- Link up the data structures
- Return a pointer to user




18

From User to System View

- What should FS do if user wants to read 10 bytes from a file starting at byte 2?
- What should FS do if user wants to write 10 bytes to a file starting at byte 2?

19

Reading a file (system call)




Modern disk drives are addressed as large one-dimensional arrays of logical blocks

20

Roadmap

- Functionality (naming/access)
 - Data structures
 - Files/directories
 - File operations
- Performance (disk management)
 - Disk layout (I/O performance)

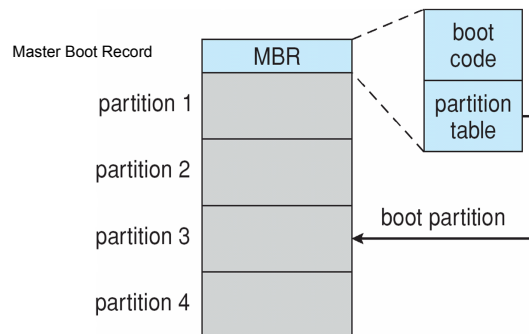


21

Boot block extra level of indirection!

- The bootstrap program is stored in ROM
 - Chip needs no initialization
 - At fixed location
 - RO means never affected by virus
 - But bootstrap program cannot be changed
 - location of OS not flexible
- Instead, modern systems store a tiny bootloader in ROM
 - which loads a full bootstrap program stored at the **boot block** at a fixed location on disk
- The full bootstrap loads the OS from non-fixed location

Booting from a Disk in Windows 2000



A Disk Layout for A File System

Boot block	Super block	File metadata (i-node in Unix)	File data blocks
------------	-------------	--------------------------------	------------------

- Boot block: contains info to boot OS
- Superblock defines a file system
 - Size of the file system
 - Free metadata (inode) count and pointers
 - Free block count and pointers (or pointer to bitmap)
 - Location of the metadata of the root directory
- What if the superblock is corrupted?
 - What can we do?

24

Disk management: Data Structures (to keep track of)



- Used space on disk:
 - A “header” for each file (part of the file meta-data)
 - Point to Disk blocks associated with each file
- Free space on disk
 - Bitmap
 - 1 bit per block (sector)
 - blocks numbered in cylinder-major order (why?)
 - Linked list of free blocks
 - How much space does a bitmap need for a 4GB disk?
 - 4294967296 bytes → 8388608 sectors → 1MB bitmap

25

Lab4 – Part 1 – FS initialization and shutdown



- We provide you the raw disk I/O
 - Simulated using the real UNIX file system
 - DiskCreate()
 - DiskWriteBlock()
 - DiskReadBlock()
- You are asked to implement
 - NewFileSys()

Boot block	Super block	File metadata (i-node in Unix)	File data blocks
---------------	----------------	-----------------------------------	------------------

26

Lab4 – Part 2 – low-level FS operations



- OpenFileSys
- CloseFileSys()
- AllocateBlock(): allocate a file system block
- FreeBlock(): free a file system block
- ReadBlock(): read a file system block
- WriteBlock(): write a file system block
- ...

27

Lab3 – Part 3 – inode based functions



- Dfs_inodeOpen()
- Dfs_inodeDelete()
- Dfs_inodeReadBytes()
- Dfs_inodeWriteBytes()
- Dfs_inodeAllocateVirtualBlock(handle,
virtual_blocknum)
- Dfs_inodeTranslateVirtualtoFilesus(handle,
virtual_blocknum)

28

Lab4 – Part 4 – basic file operations



- file_open()
- file_read()
- file_write()
- file_close()
- file_seek()
- file_chmod()
- file_delete()
- ...

How to deal with concurrent accesses?

29

Lab5 – Part 5 – Multi-level directory or Buffer Cache



30

Reading



- Chapters 10-11

31