

## File System Interface -- “A Tale of Two OS Components”

ECE595  
Nov 20

Y. Charlie Hu



## System Calls to UNIX File Systems

- 19 system calls into 6 categories:

Return file desp.	Assign inodes	Set file attr.	Process input/ output	Change file system	Modify view of file system
open close creat pipe dup	creat link unlink	chown chmod stat fstat	read write lseek	mount umount	chdir chroot


4

## Roadmap

- Functionality (API)
  - Basic functionality
    - Disk layout
    - File operations (open, read, write, close)
  - Directories
- Performance
  - Disk allocation
  - Buffer cache
  - File System interface
  - Disk scheduling
- Reliability
  - FS level
  - Disk level: RAID

5

## [lec20] Reading A Block



Modern disk drives are addressed as large one-dimensional arrays of logical blocks

6

## File System Interface

- How do application programs typically access file data?
  - Explicit read/write operations (conventional)

7

## Read / Write Interface

- File data is explicitly copied between disk file and process memory
- Programs cannot directly access file data


```
FileDescriptor fhandle;
int offset, length;
char buffer[1024];
```

```
fhandle = open("pathname");
read(fhandle, offset, buffer, length);
{read file data in buffer to do important computation};
Write(fhandle, offset, buffer, length)
close(fhandle);
```

8


## Read / Write Interface – Problem 1

- Potential for **double copies** in mem



9


## Essence of Demand Paging



Disk is the backing store

10


## [lec14] Demand Paging, Page Fault Handling



11


## Read / Write Interface – Problem 2

- Potential for **double paging** on disk
  - process pages containing file data are paged out to paging space, leading to redundant copies of file data on disk



12


## Essence of Demand Paging



Disk is the backing store

13

## What if



14

## Memory-mapped Files

- File is “mapped” into application’s address space
  - by initializing virtual memory so that the file (directly) serves as backing store for a region of the application’s address space

15

## Memory-mapped Files (cont’)

- Elegant integration of file system and virtual memory

```
FileDescriptor fhandle;  
int offset, length;  
char *address;  
  
fhandle = open("pathname");  
mmap(fhandle, offset, address, length);  
{read/write file data by accessing memory range  
  [address, address + length]};  
munmap(address, length);  
close(fhandle);
```

This is like after  
address = malloc()

16

## Memory-mapped Files

- File is “mapped” into application’s address space
  - by initializing virtual memory so that the file (directly) serves as backing store for a region of the application’s address space
- File data is *demand paged* upon access to the mapped file
  - No double paging on disk
- Memory-mapped files do not go through buffer cache
  - No double copy in mem
    - Program accesses file data directly

17

## Effects and Semantics of Memory-mapped Files

- Processes that map the same file share **physical memory** that caches file data
- Writes may not be immediately written to the file on disk
  - Update periodically
  - Closing the file results in writing all to disk and removing the VM mapping

18

## File caching implementation – Approach 1



- Set of kernel buffers maintained by the file system (buffer cache)
  - With a read/write API, can implement precise LRU
  - Not used for memory-mapped files
- Need to decide how to partition physical memory among buffer cache and VM cache uses
  - Static partitioning during kernel configuration (BSD)
  - Dynamic adjustment of partitioning during runtime, e.g. keep miss frequencies of VM and buffer cache, and try to balance them (Linux)

19

## File caching implementation – Approach 2



- File system memory-map all open files
  - Caching comes for free (just like normal VM)
  - No separate file cache
  - Flexible sharing of physical memory
  - VM page replacement policy does not discriminate between cached file data, and other VM pages
- Open/read/write API can be supported by
  - `open()` → mapping opened files into kernel address space (also using VM)
  - `read()` / `write()` → copying from/to user buffers

20

## Fun with memory-mapped files



- Inter-process communication
  - Virtual addresses of diff processes mapped to the same file
- File copying as Memory copying
  - Map files to virtual addresses
  - Do memory copying

21

## Reading



- Chapters 11-12

22